

RocksDB Festival

RF3_Team_WAL

Supported by IITP, StarLab.

July 5, 2021

김민준, 이빈

alswnssl0528@naver.com,
32183118@dankook.ac.kr

Team Name


RocksDB Festival

■ Content

- ✓ WAL Log File Format
- ✓ Experiment : is kBlockSize affected performance?
 - Info
 - Hypothesis
 - Result
 - Discussion1
 - Discussion2
- ✓ Next assignment : WAL performance according to value or key size
- ✓ Appendix : Type

RocksDB Festival : Log File Format

■ WAL Log File Format



Rn : variable size records

- ✓ Consists of a sequence of **variable** length records.
- ✓ Records are grouped by **kBlockSize**(32k).
- ✓ If a certain record cannot fit into the leftover space (leftover < Rn), then the leftover space is **padded** with empty (null) data.
- ✓ If record is bigger than kBlockSize, record occurs **fragmentation**.

RocksDB Festival : Log File Format

■ The Legacy Record Format

CRC (4B)	Size(2B)	Type(1B)	Payload
----------	----------	----------	---------

- ✓ Record consists of CRC, Size, Type, Payload
 - CRC(Cyclic Redundancy Check) : Verifies the integrity of the WAL
 - Size : Length of the record size
 - Type : kZeroType, kFullType, kFirstType, kLastType, kMiddleType

Block	FULL	A 1000 bytes	FIRST	B 31754 bytes
Block	MIDDLE	B 32761 bytes		
Block	LAST	B 32755 bytes		
Block	FULL	C 8000 bytes	6 bytes 0	

知乎 @zw Huang

- Payload : The actual value of the key-value is written

References : <https://zhuanlan.zhihu.com/p/258091002>

RocksDB Festival : Log File Format

■ The Recyclable Record Format

CRC (4B)	Size(2B)	Type(1B)	Log number (4B)	Payload
----------	----------	----------	-----------------	---------

- ✓ Record consists of CRC, Size, Type, Log Number, Payload
 - CRC, Size, Payload : same as the components of the legacy record format.
 - Type : kRecyclableFullType, kRecyclableFirstType, kRecyclableMiddleType, kRecyclableLastType
 - Log Number : Distinguish between the previous log writer and the last one. (32bit)

RocksDB Festival : kBlockSize

■ Hardware Environment : D

D	
CPU	1 * AMD Ryzen 5 3500X 6-Core
OS	Ubuntu 20.04.2 LTS
SSD	mx500

RocksDB Festival : kBlockSize

■ Experiment Info.

- ✓ WAL overhead measurement according to the kBlockSize
- ✓ Because of the kBlockSize affect the size of padding, WAL overhead will change according to the kBlockSize.

✓ Conditions


- kBlockSize = 4KB, 8KB, 16KB, 32KB(default), 64KB

[db_bench Option]


- benchmarks = "fillseq" , "fillrandom"
- disable_wal = false, true
- value_size = 16byte, 32byte, 64byte, 128byte, 256byte, 512byte, 1024byte, 2048byte, 3072byte, 4096byte, 5120byte, 6144byte, 7168byte 8192byte 16384byte

RocksDB Festival : kBlockSize

- Hypothesis - If record size is bigger than kBlockSize
 - ✓ If kBlockSize = 16KB, record size = 17KB, num=100




- ✓ Padding is **not exist**, predict performance improve.
- ✓ But, Wouldn't fragmentation cause **consistency** issues?



- ✓ Padding is exist, overhead is too big.

RocksDB Festival : kBlockSize


- Hypothesis - If kBlockSize is bigger than record size
 - ✓ Size of Padding = $kBlockSize \% \text{Size of Record}(\text{fixed})$
 - ✓ Ex. Size of record = 7KB, num=20
 - kBlockSize = 32KB, Size of Padding = 4KB
→ Total 160KB = 140KB + **20KB** (Higher overhead)
 - kBlockSize = 35KB, Size of Padding = 0KB
→ Total 140KB = 140KB + **0KB** (Lower overhead)



RocksDB Festival : kBlockSize

■ Hypothesis - Extreme situations


- ✓ If kBlockSize = 16KB, record size = 9KB, num=100



- ✓ Padding is extremely high size
 - Expected performance degradation
- ✓ Write in DB = 900KB, Write in Storage = 1600KB
 - Write Amplification is so high

RocksDB Festival : kBlockSize

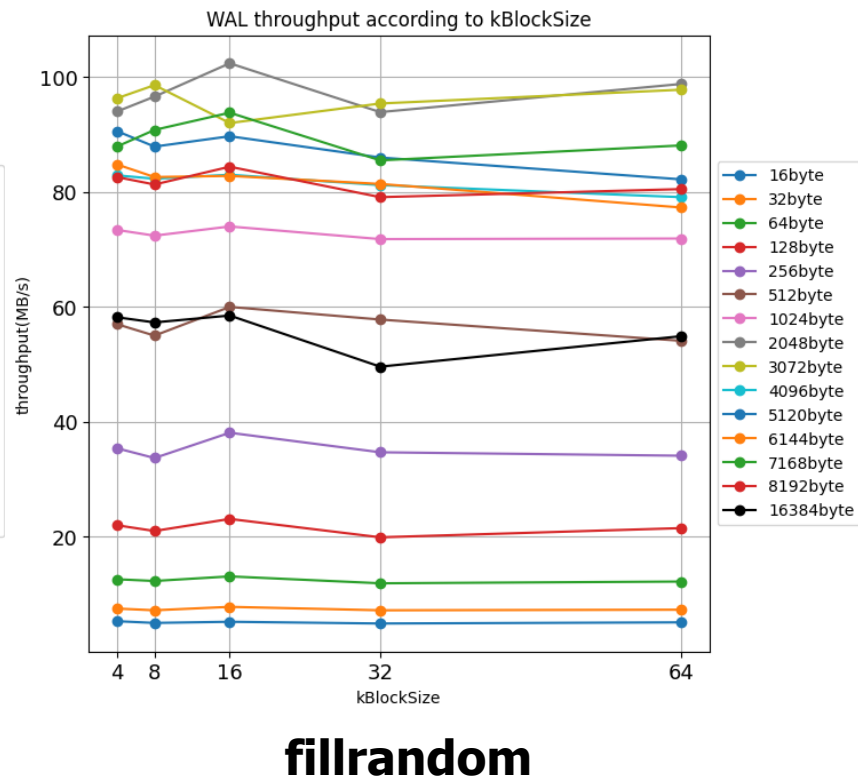

- Method for decreasing padding size
 - ✓ Predict payload size
 - ✓ Determine kBlockSize considering payload and padding size
- Despite of disadvantage, Why **kBlockSize** is used in Log File?
 - ✓ For delayed write in OS, managing static size is efficiently for processing (buffer cache)
 - ✓ Like Paging!!



RocksDB Festival : kBlockSize

■ Result 1

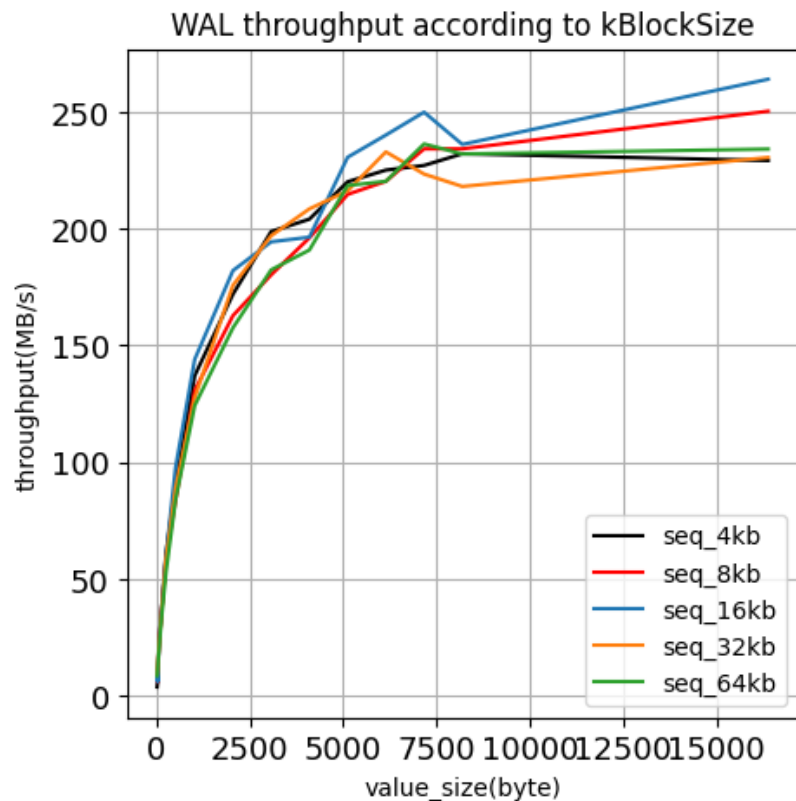
- ✓ kBlockSizes do not affect WAL overhead




RocksDB Festival : kBlockSize

■ Result 2

- ✓ kBlockSizes do not affect WAL overhead



fillseq






fillrandom

RocksDB Festival : kBlockSize

■ Discussion1

- ✓ The sum of padding sizes is constant.
- ✓ Size of record = 7KB



RocksDB Festival : kBlockSize

■ Discussion2

- ✓ Padding is **not exist!**
- ✓ Additional experiments to observe padding

```
// is empty, we still want to iterate once to emit a single
// zero-length record
IOStatus s;
bool begin = true;
do {
    const int64_t leftover = kBlockSize - block_offset_;

    if (printf(stdout, "leftover : %ld\n", leftover);


    assert(leftover >= 0);
    if (leftover < header_size) {
        // Switch to a new block
        if (leftover > 0) {
            // Fill the trailer (literal below relies on kHeaderSize and
            // kRecyclableHeaderSize being <= 11)
            assert(header_size <= 11);
            s = dest_->Append(Slice("\x00\x00\x00\x00\x00\x00\x00\x00\x00\x00'
                                     static_cast<size_t>(leftover)));
```

log_writer.cc

RocksDB Festival : kBlockSize

■ Discussion2


- ✓ Kblocksize : 8KB, Key Size : 16byte, Value Size : 1500byte
- ✓ leftover = kBlockSize - block_offset_;




RocksDB Festival : kBlockSize

■ Discussion2 - Is padding existed in kBlock?

✓ Expected



✓ Actual result



RocksDB Festival : kBlockSize

■ Source code analysis

✓ AddRecord => log file writer

```
IOStatus Writer::AddRecord(const Slice& slice) {
```

```
    const char* ptr = slice.data();
```

```
    size_t left = slice.size(); Left : record size
```

```
    // Header size varies depending on whether we are recycling or not.
```

```
    const int header_size =
```

```
        recycle_log_files_ ? kRecyclableHeaderSize : kHeaderSize;
```

```
    // Fragment the record if necessary and emit it. Note that if slice
```

```
    // is empty, we still want to iterate once to emit a single
```

```
    // zero-length record
```

```
    IOStatus s;
```

```
    bool begin = true;
```

```
    do {
```

```
        const int64_t leftover = kBlockSize - block_offset_;
```

Leftover : empty space of kBlock

```
        assert(leftover >= 0);
```

```
        if (leftover < header_size) {
```

```
            // Switch to a new block
```

```
            if (leftover > 0) {
```

```
                // Fill the trailer (literal below relies on kHeaderSize and
```

```
                // kRecyclableHeaderSize being <= 11)
```

```
                assert(header_size <= 11);
```

```
                s = dest_>Append(Slice("\x00\x00\x00\x00\x00\x00\x00\x00\x00"
                                         static_cast<size_t>(leftover)));
```

**If leftover smaller than
headersize, padding insert in
kBlock Switch to next kBlock**

```
                if (!s.ok()) {
```

```
                    break;
```

```
                }
```

```
            }
```

```
            block_offset_ = 0;
```

```
        }
```

RocksDB Festival : kBlockSize

■ Source code analysis

```
// Invariant: we never leave < header_size bytes in a block.  
assert(static_cast<int64_t>(kBlockSize - block_offset_) >= header_size);
```

```
const size_t avail = kBlockSize - block_offset_ - header_size;  
const size_t fragment_length = (left < avail) ? left : avail;
```

If left bigger than avail, it is fragmented by avail

```
RecordType type;  
const bool end = (left == fragment_length);  
if (begin && end) {  
    type = recycle_log_files_ ? kRecyclableFullType : kFullType;  
} else if (begin) {  
    type = recycle_log_files_ ? kRecyclableFirstType : kFirstType;  
} else if (end) {  
    type = recycle_log_files_ ? kRecyclableLastType : kLastType;  
} else {  
    type = recycle_log_files_ ? kRecyclableMiddleType : kMiddleType;  
}
```

Padding is created only when the header cannot be stored

```
s = EmitPhysicalRecord(type, ptr, fragment_length);  
ptr += fragment_length;  
left -= fragment_length;  
begin = false;  
} while (s.ok() && left > 0);
```

Remain record size

If record are entirely stored in kBlock, the loop ends


```
if (s.ok()) {  
    if (!manual_flush_) {  
        s = dest_>Flush();  
    }  
}
```

```
return s;
```

RocksDB Festival : key / value size

■ Next assignment

- ✓ WAL performance according to **value** or **key size**



CRC (4B)	Size(2B)	Type(1B)	Payload (Variable Length)
----------	----------	----------	---------------------------

- ✓ db_bench options
 - --disable_wal=[boolean]
 - --key_size=[int value]
 - --value_size=[int value]

Discussion



Appendix : Type

■ RecordType - 27061 bytes Records

```
leftover : 8192
Type : 2
leftover : 0
Type : 3
leftover : 0
Type : 3
leftover : 0
Type : 4
leftover : 5707
Type : 2
leftover : 0
Type : 3
leftover : 0
Type : 3
leftover : 0
Type : 4
leftover : 3222
Type : 2
```

```
enum RecordType {
    // Zero is reserved for preallocated files
    kZeroType = 0,
    kFullType = 1,

    // For fragments
    kFirstType = 2,
    kMiddleType = 3,
    kLastType = 4,
}
```

Block0

Record0 (0) kFirstType(2) 8192 bytes

Block1

Record0 (1) kMiddleType(3) 8192 bytes

Block2

Record0 (2) kMiddleType(3) 8192 bytes

Block3

Record0 (3) kLastType(4) 2485 bytes

Record1 (0)

Block4

Record1 (1) kMiddleType(3) ...