

RocksDB Festival

RF3_Team_WAL

Supported by IITP, StarLab.

July 5, 2021

김민준, 이빈

alswnssl0528@naver.com,

32183118@dankook.ac.kr

Team Name

RocksDB Festival

■ Contents

- ✓ 팀 구성 : **Team Info**
- ✓ RocksDB db_bench 설치 및 실험 결과 : **db_bench**
- ✓ 선택한 주제와 선택한 이유 : **WAL**
- ✓ 이번 스터디 동안 기대하는 결과

RocksDB Festival : Team Info

■ Team Member (Github ID)

- ✓ 김민준(Sadsprin)
- ✓ alswnssl0528@naver.com

- ✓ 이빈(Bynnn)
- ✓ 32183118@dankook.ac.kr

RocksDB Festival : db_bench

■ RocksDB db_bench 설치 및 실험 결과

✓ 실험 환경 : A, B, C

A

CPU	4 * Intel(R) Core(TM) i5-7300HQ CPU @ 2.50GHZ
OS	Ubuntu 20.04.2 LTS
	Real Machine

B

CPU	2 * Intel(R) Core(TM) i5-4690 CPU @ 3.50GHZ
OS	Ubuntu 20.04.2 LTS
	Virtual Machine

C

CPU	4 * AMD Ryzen 7 3700X 8-Core Processor
OS	Ubuntu 20.04.2 LTS
SSD	PM981a
	Virtual Machine

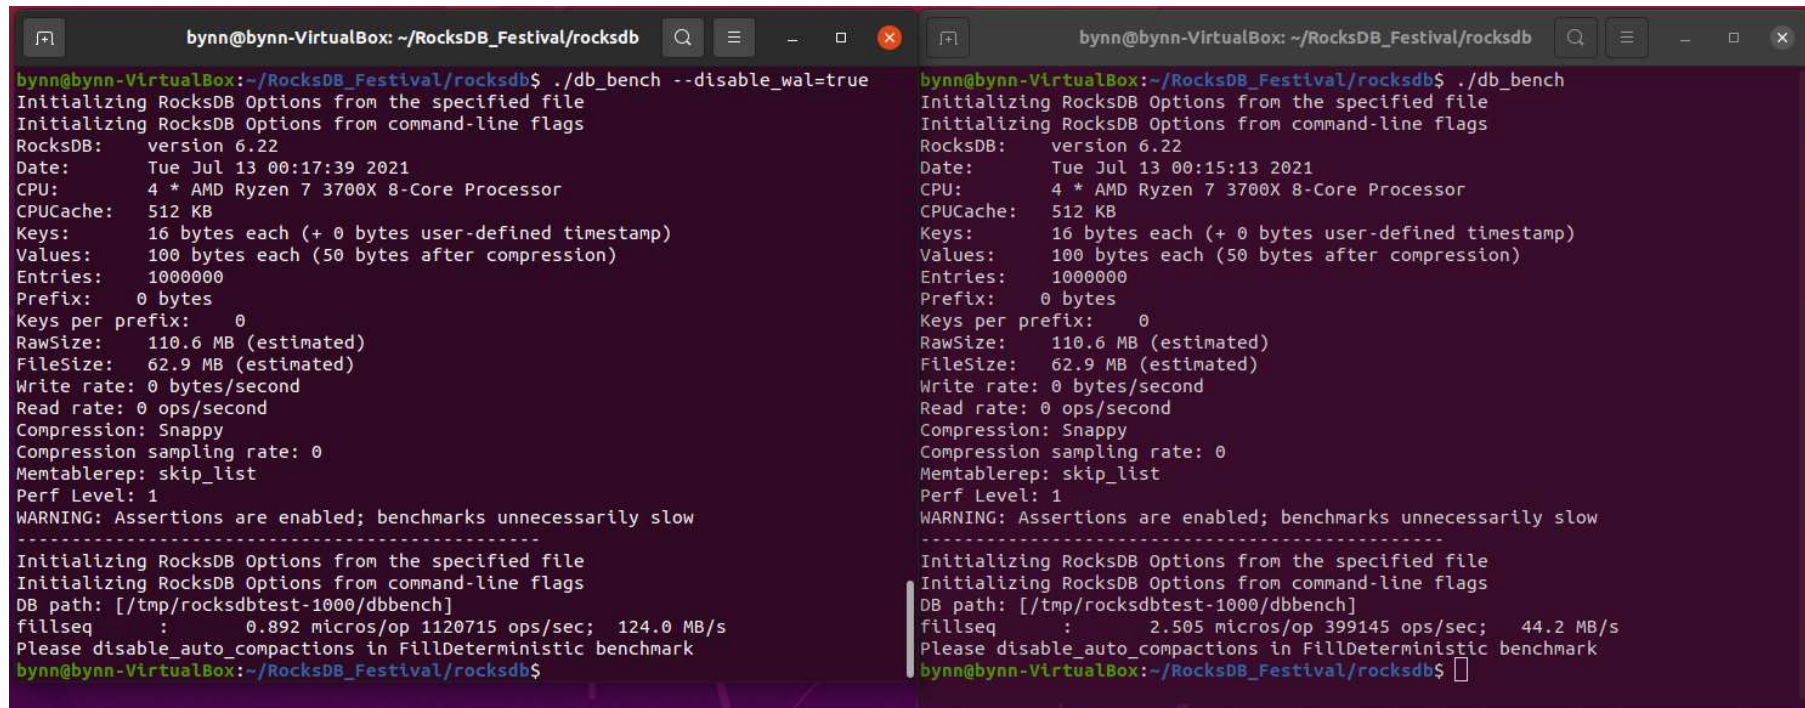
RocksDB Festival : db_bench

■ RocksDB db_bench 설치 및 실험 결과

```
bynn@bynn-VirtualBox:~/RocksDB_Festival/rocksdb$ ls
appveyor.yml          examples              options
AUTHORS               file                 plugin
buckifier             fuzz                 PLUGINS.md
build_tools           hdfs                 port
cache                 HISTORY.md           README.md
cmake                  include              ROCKSDB_LITE.md
CMakeLists.txt        INSTALL.md           src.mk
CODE_OF_CONDUCT.md    issue_template.md   table
CONTRIBUTING.md      java                 TARGETS
COPYING               LANGUAGE-BINDINGS.md test_util
coverage              librocksdb_debug.a  third-party
db                    LICENSE.Apache       thirdparty.inc
db_bench              LICENSE.leveldb     tools
db_stress_tool        logging              trace_replay
DEFAULT_OPTIONS_HISTORY.md make_config.mk       USERS.md
defs.bzl              Makefile             util
docs                  memory               utilities
DUMP_FORMAT.md        memtable             Vagrantfile
env                   monitoring           WINDOWS_PORT.md
```

RocksDB Festival : db_bench

- RocksDB db_bench 설치 및 실험 결과
 - ✓ --disable_wal=[boolean]



```
bynn@bynn-VirtualBox: ~/RocksDB_Festival/rocksdb$ ./db_bench --disable_wal=true
Initializing RocksDB Options from the specified file
Initializing RocksDB Options from command-line flags
RocksDB: version 6.22
Date: Tue Jul 13 00:17:39 2021
CPU: 4 * AMD Ryzen 7 3700X 8-Core Processor
CPU Cache: 512 KB
Keys: 16 bytes each (+ 0 bytes user-defined timestamp)
Values: 100 bytes each (50 bytes after compression)
Entries: 1000000
Prefix: 0 bytes
Keys per prefix: 0
RawSize: 110.6 MB (estimated)
FileSize: 62.9 MB (estimated)
Write rate: 0 bytes/second
Read rate: 0 ops/second
Compression: Snappy
Compression sampling rate: 0
Memtablerep: skip_list
Perf Level: 1
WARNING: Assertions are enabled; benchmarks unnecessarily slow
-----
Initializing RocksDB Options from the specified file
Initializing RocksDB Options from command-line flags
DB path: [/tmp/rocksdbtest-1000/dbbench]
fillseq : 0.892 micros/op 1120715 ops/sec; 124.0 MB/s
Please disable_auto_compactions in FillDeterministic benchmark
bynn@bynn-VirtualBox:~/RocksDB_Festival/rocksdb$

bynn@bynn-VirtualBox:~/RocksDB_Festival/rocksdb$ ./db_bench
Initializing RocksDB Options from the specified file
Initializing RocksDB Options from command-line flags
RocksDB: version 6.22
Date: Tue Jul 13 00:15:13 2021
CPU: 4 * AMD Ryzen 7 3700X 8-Core Processor
CPU Cache: 512 KB
Keys: 16 bytes each (+ 0 bytes user-defined timestamp)
Values: 100 bytes each (50 bytes after compression)
Entries: 1000000
Prefix: 0 bytes
Keys per prefix: 0
RawSize: 110.6 MB (estimated)
FileSize: 62.9 MB (estimated)
Write rate: 0 bytes/second
Read rate: 0 ops/second
Compression: Snappy
Compression sampling rate: 0
Memtablerep: skip_list
Perf Level: 1
WARNING: Assertions are enabled; benchmarks unnecessarily slow
-----
Initializing RocksDB Options from the specified file
Initializing RocksDB Options from command-line flags
DB path: [/tmp/rocksdbtest-1000/dbbench]
fillseq : 2.505 micros/op 399145 ops/sec; 44.2 MB/s
Please disable_auto_compactions in FillDeterministic benchmark
bynn@bynn-VirtualBox:~/RocksDB_Festival/rocksdb$
```


RocksDB Festival : db_bench

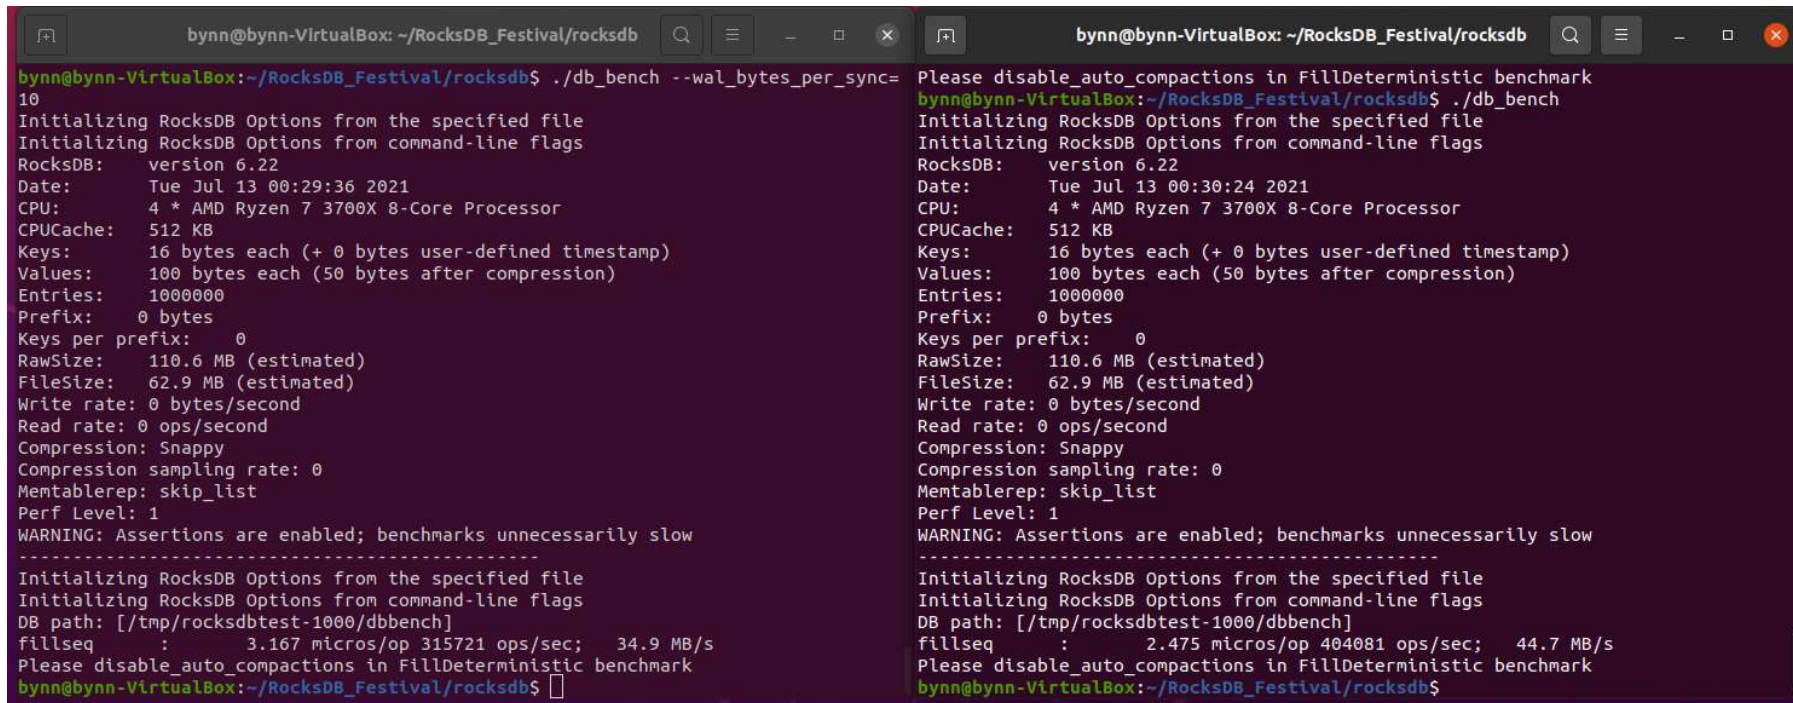
- RocksDB db_bench 설치 및 실험 결과
 - ✓ --max_total_wal_size=[int value]

```
bynn@bynn-VirtualBox: ~/RocksDB_Festival/rocksdb
bynn@bynn-VirtualBox:~/RocksDB_Festival/rocksdb$ ./db_bench --max_total_wal_size=10
Please disable_auto_compactions in FillDeterministic benchmark
bynn@bynn-VirtualBox:~/RocksDB_Festival/rocksdb$ ./db_bench
Initializing RocksDB Options from the specified file
Initializing RocksDB Options from command-line flags
RocksDB: version 6.22
Date: Tue Jul 13 00:26:05 2021
CPU: 4 * AMD Ryzen 7 3700X 8-Core Processor
CPUCache: 512 KB
Keys: 16 bytes each (+ 0 bytes user-defined timestamp)
Values: 100 bytes each (50 bytes after compression)
Entries: 1000000
Prefix: 0 bytes
Keys per prefix: 0
RawSize: 110.6 MB (estimated)
FileSize: 62.9 MB (estimated)
Write rate: 0 bytes/second
Read rate: 0 ops/second
Compression: Snappy
Compression sampling rate: 0
Memtablerep: skip_list
Perf Level: 1
WARNING: Assertions are enabled; benchmarks unnecessarily slow
-----
Initializing RocksDB Options from the specified file
Initializing RocksDB Options from command-line flags
DB path: [/tmp/rocksdbtest-1000/dbbench]
fillseq : 3.938 micros/op 253951 ops/sec; 28.1 MB/s
Please disable_auto_compactions in FillDeterministic benchmark
bynn@bynn-VirtualBox:~/RocksDB_Festival/rocksdb$

bynn@bynn-VirtualBox:~/RocksDB_Festival/rocksdb$ ./db_bench
Please disable_auto_compactions in FillDeterministic benchmark
bynn@bynn-VirtualBox:~/RocksDB_Festival/rocksdb$ ./db_bench
Initializing RocksDB Options from the specified file
Initializing RocksDB Options from command-line flags
RocksDB: version 6.22
Date: Tue Jul 13 00:25:30 2021
CPU: 4 * AMD Ryzen 7 3700X 8-Core Processor
CPUCache: 512 KB
Keys: 16 bytes each (+ 0 bytes user-defined timestamp)
Values: 100 bytes each (50 bytes after compression)
Entries: 1000000
Prefix: 0 bytes
Keys per prefix: 0
RawSize: 110.6 MB (estimated)
FileSize: 62.9 MB (estimated)
Write rate: 0 bytes/second
Read rate: 0 ops/second
Compression: Snappy
Compression sampling rate: 0
Memtablerep: skip_list
Perf Level: 1
WARNING: Assertions are enabled; benchmarks unnecessarily slow
-----
Initializing RocksDB Options from the specified file
Initializing RocksDB Options from command-line flags
DB path: [/tmp/rocksdbtest-1000/dbbench]
fillseq : 2.474 micros/op 404127 ops/sec; 44.7 MB/s
Please disable_auto_compactions in FillDeterministic benchmark
bynn@bynn-VirtualBox:~/RocksDB_Festival/rocksdb$
```

RocksDB Festival : db_bench

- RocksDB db_bench 설치 및 실험 결과
 - ✓ --wal_bytes_per_sync=[int value]

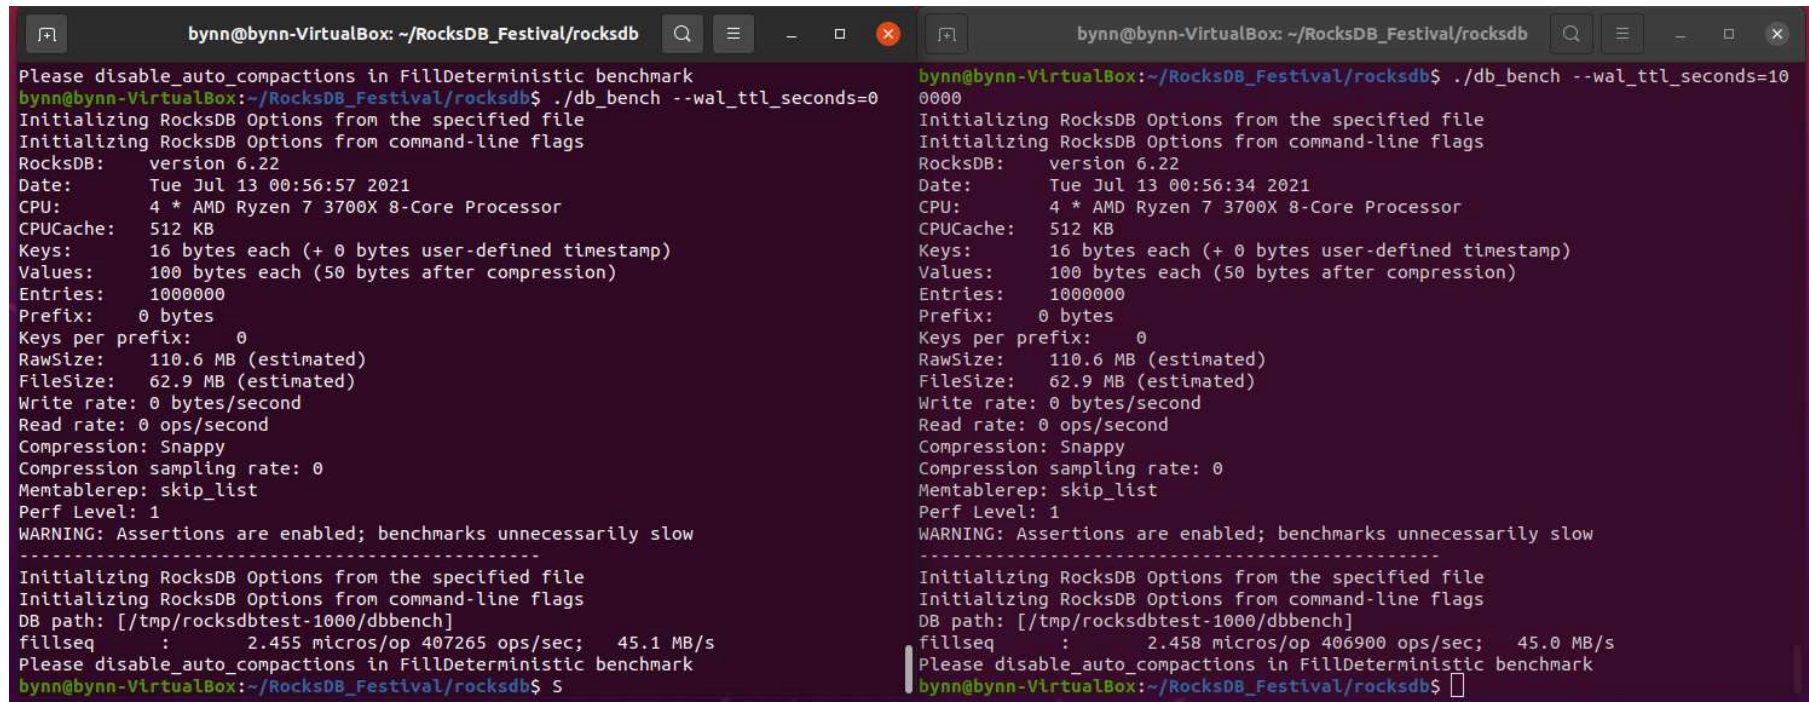


```
bynn@bynn-VirtualBox: ~/RocksDB_Festival/rocksdb$ ./db_bench --wal_bytes_per_sync=10
Initializing RocksDB Options from the specified file
Initializing RocksDB Options from command-line flags
RocksDB: version 6.22
Date: Tue Jul 13 00:29:36 2021
CPU: 4 * AMD Ryzen 7 3700X 8-Core Processor
CPUCache: 512 KB
Keys: 16 bytes each (+ 0 bytes user-defined timestamp)
Values: 100 bytes each (50 bytes after compression)
Entries: 1000000
Prefix: 0 bytes
Keys per prefix: 0
RawSize: 110.6 MB (estimated)
FileSize: 62.9 MB (estimated)
Write rate: 0 bytes/second
Read rate: 0 ops/second
Compression: Snappy
Compression sampling rate: 0
Memtablerep: skip_list
Perf Level: 1
WARNING: Assertions are enabled; benchmarks unnecessarily slow
-----
Initializing RocksDB Options from the specified file
Initializing RocksDB Options from command-line flags
DB path: [/tmp/rocksdbtest-1000/dbbench]
fillseq : 3.167 micros/op 315721 ops/sec; 34.9 MB/s
Please disable_auto_compactions in FillDeterministic benchmark
bynn@bynn-VirtualBox:~/RocksDB_Festival/rocksdb$

bynn@bynn-VirtualBox: ~/RocksDB_Festival/rocksdb$ ./db_bench
Please disable_auto_compactions in FillDeterministic benchmark
bynn@bynn-VirtualBox:~/RocksDB_Festival/rocksdb$ ./db_bench
Initializing RocksDB Options from the specified file
Initializing RocksDB Options from command-line flags
RocksDB: version 6.22
Date: Tue Jul 13 00:30:24 2021
CPU: 4 * AMD Ryzen 7 3700X 8-Core Processor
CPUCache: 512 KB
Keys: 16 bytes each (+ 0 bytes user-defined timestamp)
Values: 100 bytes each (50 bytes after compression)
Entries: 1000000
Prefix: 0 bytes
Keys per prefix: 0
RawSize: 110.6 MB (estimated)
FileSize: 62.9 MB (estimated)
Write rate: 0 bytes/second
Read rate: 0 ops/second
Compression: Snappy
Compression sampling rate: 0
Memtablerep: skip_list
Perf Level: 1
WARNING: Assertions are enabled; benchmarks unnecessarily slow
-----
Initializing RocksDB Options from the specified file
Initializing RocksDB Options from command-line flags
DB path: [/tmp/rocksdbtest-1000/dbbench]
fillseq : 2.475 micros/op 404081 ops/sec; 44.7 MB/s
Please disable_auto_compactions in FillDeterministic benchmark
bynn@bynn-VirtualBox:~/RocksDB_Festival/rocksdb$
```


RocksDB Festival : db_bench

- RocksDB db_bench 설치 및 실험 결과
 - ✓ --wal_ttl_seconds=[int value]



```
bynn@bynn-VirtualBox: ~/RocksDB_Festival/rocksdb
Please disable_auto_compactions in FillDeterministic benchmark
bynn@bynn-VirtualBox:~/RocksDB_Festival/rocksdb$ ./db_bench --wal_ttl_seconds=0
Initializing RocksDB Options from the specified file
Initializing RocksDB Options from command-line flags
RocksDB: version 6.22
Date: Tue Jul 13 00:56:57 2021
CPU: 4 * AMD Ryzen 7 3700X 8-Core Processor
CPUCache: 512 KB
Keys: 16 bytes each (+ 0 bytes user-defined timestamp)
Values: 100 bytes each (50 bytes after compression)
Entries: 1000000
Prefix: 0 bytes
Keys per prefix: 0
RawSize: 110.6 MB (estimated)
FileSize: 62.9 MB (estimated)
Write rate: 0 bytes/second
Read rate: 0 ops/second
Compression: Snappy
Compression sampling rate: 0
Mentablerep: skip_list
Perf Level: 1
WARNING: Assertions are enabled; benchmarks unnecessarily slow
-----
Initializing RocksDB Options from the specified file
Initializing RocksDB Options from command-line flags
DB path: [/tmp/rocksdbtest-1000/dbbench]
fillseq : 2.455 micros/op 407265 ops/sec; 45.1 MB/s
Please disable_auto_compactions in FillDeterministic benchmark
bynn@bynn-VirtualBox:~/RocksDB_Festival/rocksdb$

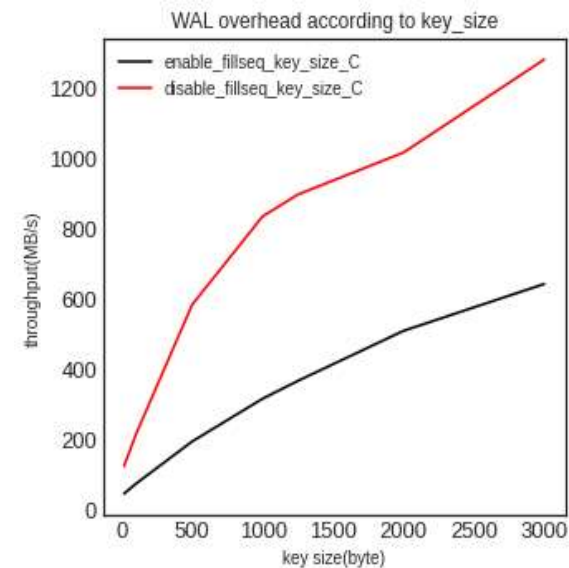
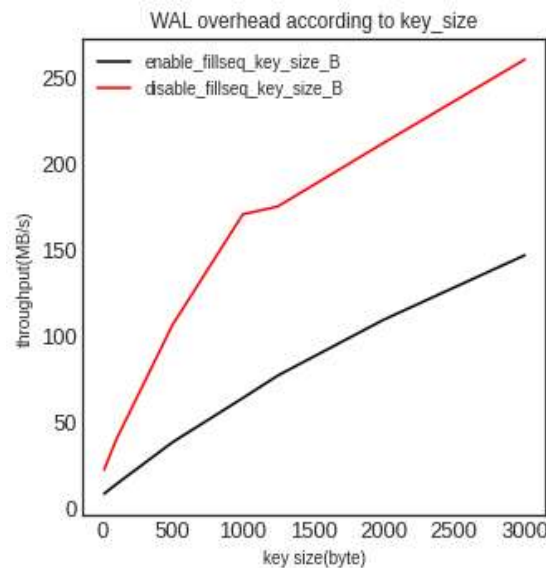
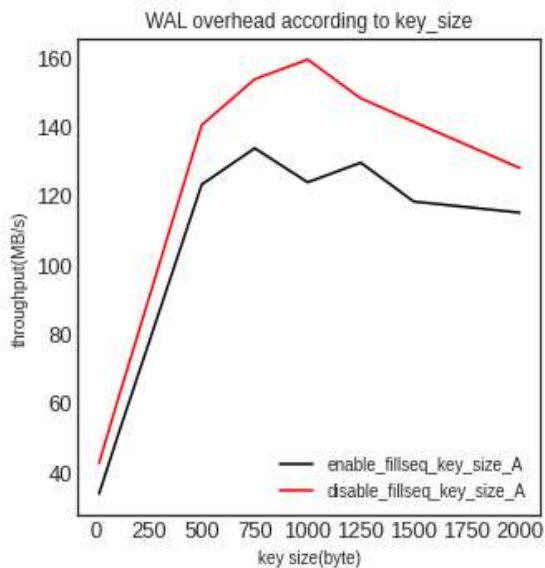
bynn@bynn-VirtualBox: ~/RocksDB_Festival/rocksdb
bynn@bynn-VirtualBox:~/RocksDB_Festival/rocksdb$ ./db_bench --wal_ttl_seconds=10
0000
Initializing RocksDB Options from the specified file
Initializing RocksDB Options from command-line flags
RocksDB: version 6.22
Date: Tue Jul 13 00:56:34 2021
CPU: 4 * AMD Ryzen 7 3700X 8-Core Processor
CPUCache: 512 KB
Keys: 16 bytes each (+ 0 bytes user-defined timestamp)
Values: 100 bytes each (50 bytes after compression)
Entries: 1000000
Prefix: 0 bytes
Keys per prefix: 0
RawSize: 110.6 MB (estimated)
FileSize: 62.9 MB (estimated)
Write rate: 0 bytes/second
Read rate: 0 ops/second
Compression: Snappy
Compression sampling rate: 0
Mentablerep: skip_list
Perf Level: 1
WARNING: Assertions are enabled; benchmarks unnecessarily slow
-----
Initializing RocksDB Options from the specified file
Initializing RocksDB Options from command-line flags
DB path: [/tmp/rocksdbtest-1000/dbbench]
fillseq : 2.458 micros/op 406900 ops/sec; 45.0 MB/s
Please disable_auto_compactions in FillDeterministic benchmark
bynn@bynn-VirtualBox:~/RocksDB_Festival/rocksdb$
```

RocksDB Festival : db_bench

■ RocksDB db_bench 설치 및 실험 결과


✓ WAL Overhead 측정 실험

- RocksDB에서 WAL 로깅에 의한 overhead를 측정한다.
- Overhead : WAL 사용 유무에 대한 Throughput(ops/sec(MB/s))의 차이
- db_bench options
 - --disable_wal=[boolean]
 - --key_size=[int value]



RocksDB Festival : WAL

- 선택한 주제와 선택한 이유
 - ✓ WAL(Write-Ahead-Log)
 - ✓ Every update to RocksDB is written to WAL on disk.
 - ✓ In the event of a failure, write ahead logs can be used to completely recover the data in the memtable, which is necessary to restore the database to the original state.
 - ✓ In the default configuration, RocksDB guarantees process crash consistency by flushing the WAL after every user write.



RocksDB Festival

- 이번 스터디 동안 기대하는 결과
 1. WAL의 구조에 따른 오버헤드 분석
 2. WAL 오버헤드를 줄이기 위한 방법 고안
 - > WAL format, total size of WAL, compression
 - > system failure, system crash prediction

Discussion

