

RocksDB Festival

RF3_Team_WAL

Supported by IITP, StarLab.

July 5, 2021

김민준, 이빈

alswnssl0528@naver.com,
32183118@dankook.ac.kr

Team Name


RocksDB Festival

■ Content

- ✓ WAL Log File Format
- ✓ Experiment : is kBlockSize affected performance?
 - Info
 - Hypothesis
 - Result
 - Discussion1
 - Discussion2
- ✓ Next assignment : WAL performance according to value or key size

RocksDB Festival : Log File Format

■ WAL Log File Format



Rn : variable size records

- ✓ Consists of a sequence of **variable** length records.
- ✓ Records are grouped by **kBlockSize**(32k).
- ✓ If a certain record cannot fit into the leftover space (leftover < Rn), then the leftover space is **padded** with empty (null) data.
- ✓ If record is bigger than kBlockSize, record occurs **fragmentation**.

RocksDB Festival : Log File Format

■ The Legacy Record Format

CRC (4B)	Size(2B)	Type(1B)	Payload
----------	----------	----------	---------

- ✓ Record consists of CRC, Size, Type, Payload
 - CRC(Cyclic Redundancy Check) : Verifies the integrity of the WAL
 - Size : Length of the record size
 - Type : kZeroType, kFullType, kFirstType, kLastType, kMiddleType

Block	FULL	A 1000 bytes	FIRST	B 31754 bytes
Block	MIDDLE	B 32761 bytes		
Block	LAST	B 32755 bytes		
Block	FULL	C 8000 bytes		6 bytes 0

知乎 @zw Huang

- Payload : The actual value of the key-value is written

References : <https://zhuanlan.zhihu.com/p/258091002>

RocksDB Festival : Log File Format

■ The Recyclable Record Format

CRC (4B)	Size(2B)	Type(1B)	Log number (4B)	Payload
----------	----------	----------	-----------------	---------

- ✓ Record consists of CRC, Size, Type, Log Number, Payload
 - CRC, Size, Payload : same as the components of the legacy record format.
 - Type : kRecyclableFullType, kRecyclableFirstType, kRecyclableMiddleType, kRecyclableLastType
 - Log Number : Distinguish between the previous log writer and the last one. (32bit)

RocksDB Festival : kBlockSize

■ Hardware Environment : D

D	
CPU	1 * AMD Ryzen 5 3500X 6-Core
OS	Ubuntu 20.04.2 LTS
SSD	mx500

RocksDB Festival : kBlockSize

■ Experiment Info.

- ✓ WAL overhead measurement according to the kBlockSize
- ✓ Because of the kBlockSize affect the size of padding, WAL overhead will change according to the kBlockSize.

✓ Conditions


- kBlockSize = 4KB, 8KB, 16KB, 32KB(default), 64KB

[db_bench Option]


- benchmarks = "fillseq" , "fillrandom"
- disable_wal = false, true
- value_size = 16byte, 32byte, 64byte, 128byte, 256byte, 512byte, 1024byte, 2048byte, 3072byte, 4096byte, 5120byte, 6144byte, 7168byte 8192byte 16384byte

RocksDB Festival : kBlockSize

- Hypothesis - If record size is bigger than kBlockSize
 - ✓ If kBlockSize = 16KB, record size = 17KB, num=100




- ✓ Padding is **not exist**, predict performance improve.
- ✓ But, Wouldn't fragmentation cause **consistency** issues?



- ✓ Padding is exist, overhead is too big.

RocksDB Festival : kBlockSize


- Hypothesis - If kBlockSize is bigger than record size
 - ✓ Size of Padding = $kBlockSize \% \text{Size of Record}(\text{fixed})$
 - ✓ Ex. Size of record = 7KB, num=20
 - kBlockSize = 32KB, Size of Padding = 4KB
→ Total 160KB = 140KB + **20KB** (Higher overhead)
 - kBlockSize = 35KB, Size of Padding = 0KB
→ Total 140KB = 140KB + **0KB** (Lower overhead)



RocksDB Festival : kBlockSize

■ Hypothesis - Extreme situations


- ✓ If kBlockSize = 16KB, record size = 9KB, num=100



- ✓ Padding is extremely high size
 - Expected performance degradation
- ✓ Write in DB = 900KB, Write in Storage = 1600KB
 - Write Amplification is so high

RocksDB Festival : kBlockSize



- Method for decreasing padding size
 - ✓ Predict payload size
 - ✓ Determine kBlockSize considering payload and padding size
- Despite of disadvantage, Why **kBlockSize** is used in Log File?
 - ✓ For delayed write in OS, managing static size is efficiently for processing (buffer cache)
 - ✓ Like Paging!!



RocksDB Festival : kBlockSize

■ Result 1

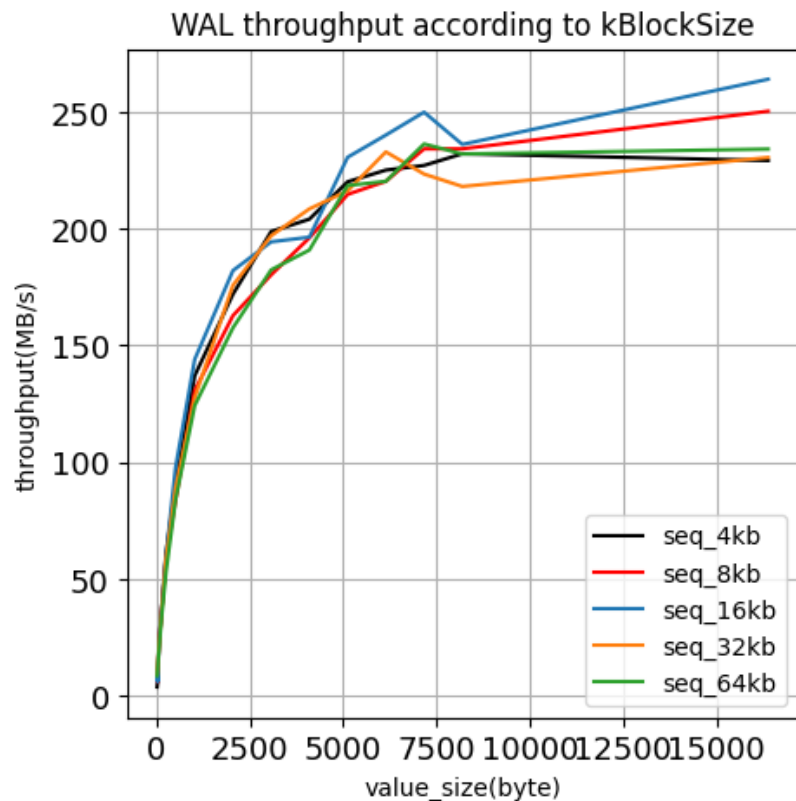
- ✓ kBlockSize do not affect WAL overhead




RocksDB Festival : kBlockSize

■ Result 2

- ✓ kBlockSizes do not affect WAL overhead



fillseq






fillrandom

RocksDB Festival : kBlockSize

■ Discussion1

- ✓ The sum of padding sizes is constant.
- ✓ Size of record = 7KB



RocksDB Festival : kBlockSize

■ Discussion2

- ✓ Padding is **not exist!**
- ✓ Additional experiments to observe padding

```
// is empty, we still want to iterate once to emit a single
// zero-length record
IOStatus s;
bool begin = true;
do {
    const int64_t leftover = kBlockSize - block_offset_;

    if (printf(stdout, "leftover : %ld\n", leftover);


    assert(leftover >= 0);
    if (leftover < header_size) {
        // Switch to a new block
        if (leftover > 0) {
            // Fill the trailer (literal below relies on kHeaderSize and
            // kRecyclableHeaderSize being <= 11)
            assert(header_size <= 11);
            s = dest_->Append(Slice("\x00\x00\x00\x00\x00\x00\x00\x00\x00\x00'
                                     static_cast<size_t>(leftover)));
```

log_writer.cc

RocksDB Festival : kBlockSize

■ Discussion2


- ✓ Kblocksize : 8KB, Key Size : 16byte, Value Size : 1500byte
- ✓ leftover = kBlockSize - block_offset_;




RocksDB Festival : kBlockSize

■ Discussion2 - Is padding existed in kBlock?

✓ Expected




✓ Actual result



RocksDB Festival : key / value size

■ Next assignment

- ✓ WAL performance according to **value** or **key size**



CRC (4B)	Size(2B)	Type(1B)	Payload (Variable Length)
----------	----------	----------	---------------------------

- ✓ db_bench options
 - --disable_wal=[boolean]
 - --key_size=[int value]
 - --value_size=[int value]

Discussion

