



RocksDB Festival

RocksDB Architecture and Operation

Supported by IITP, StarLab.

July 12, 2021

Hojin Shin, Jongmoo Choi

choijm@dankook.ac.kr

<http://embedded.dankook.ac.kr/~choijm>

RocksDB Festival: Contents

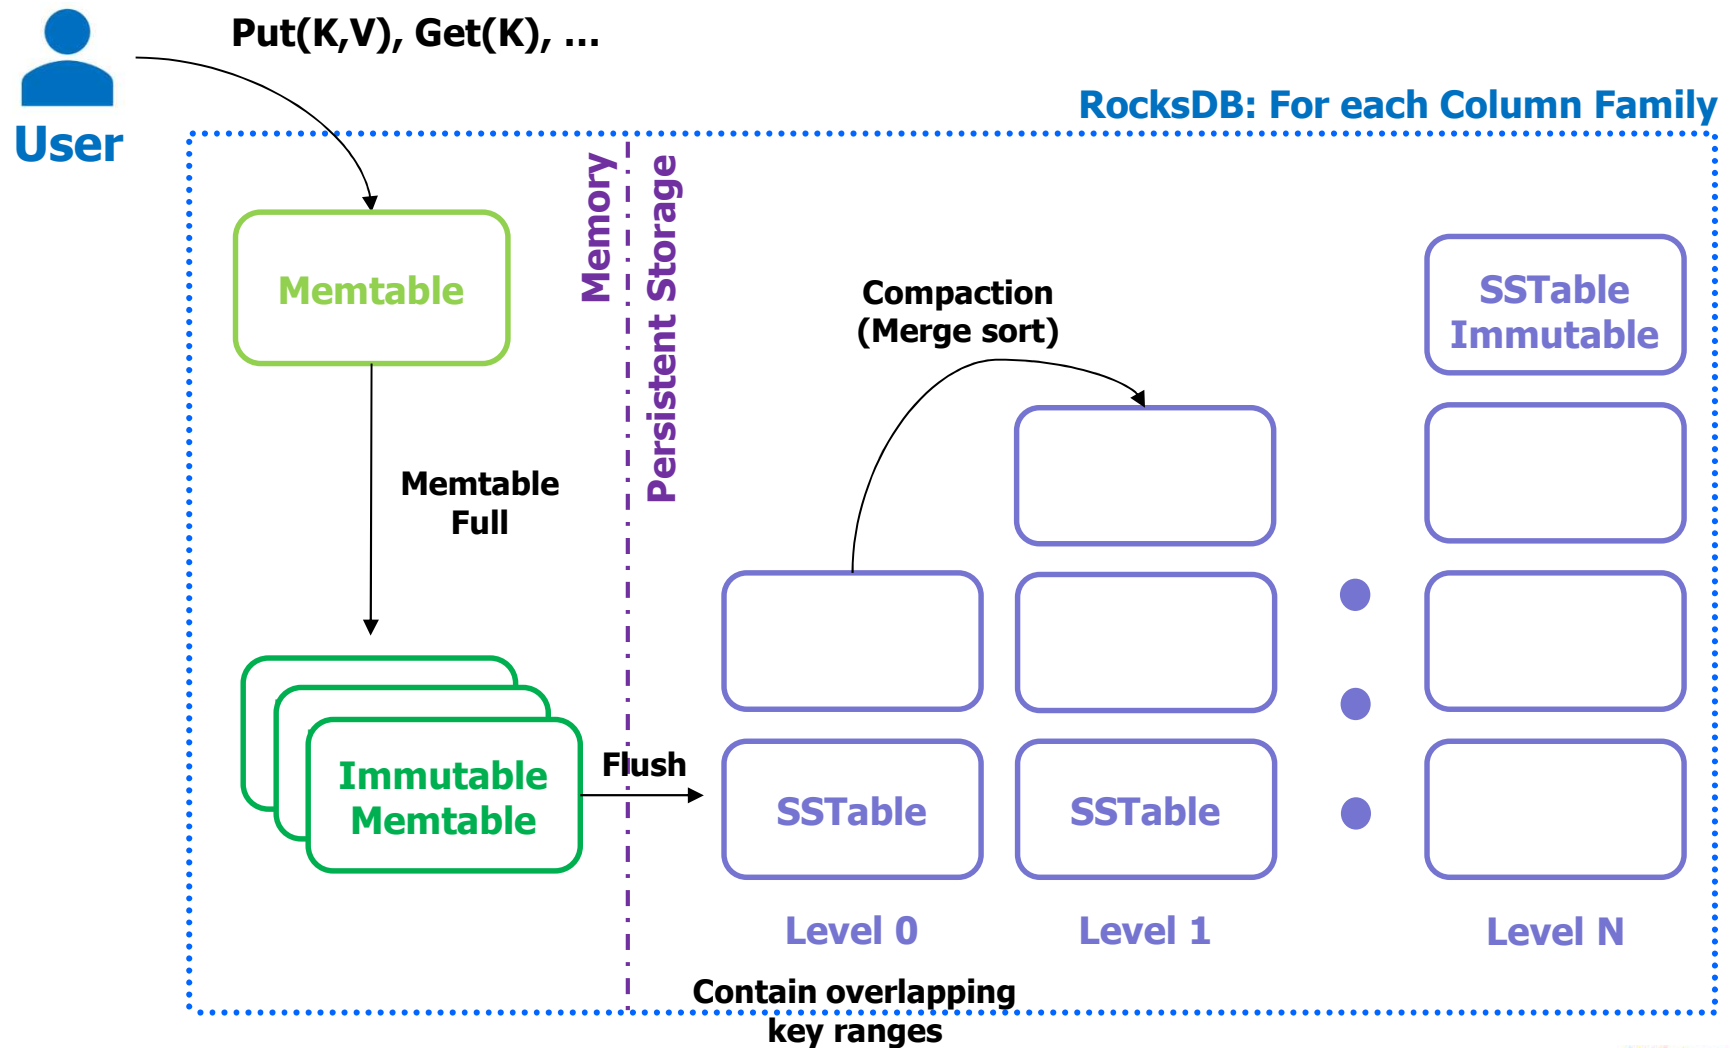
■ Contents

- ✓ RocksDB Architecture: Overall
- ✓ RocksDB Architecture: MemTable
- ✓ RocksDB Architecture: SSTable
- ✓ RocksDB Operation: Write and WAL
- ✓ RocksDB Operation: Read
- ✓ RocksDB Operation: Compaction
- ✓ QnA




RocksDB Festival: Overall

■ RocksDB Architecture



RocksDB Festival: Overall

■ RocksDB Architecture: Vending Machine



**Get(K)
??**

Vending Machine Database

Key	Value
초코우유1	700
초코우유2	700
딸기우유	700
흰 우유	500

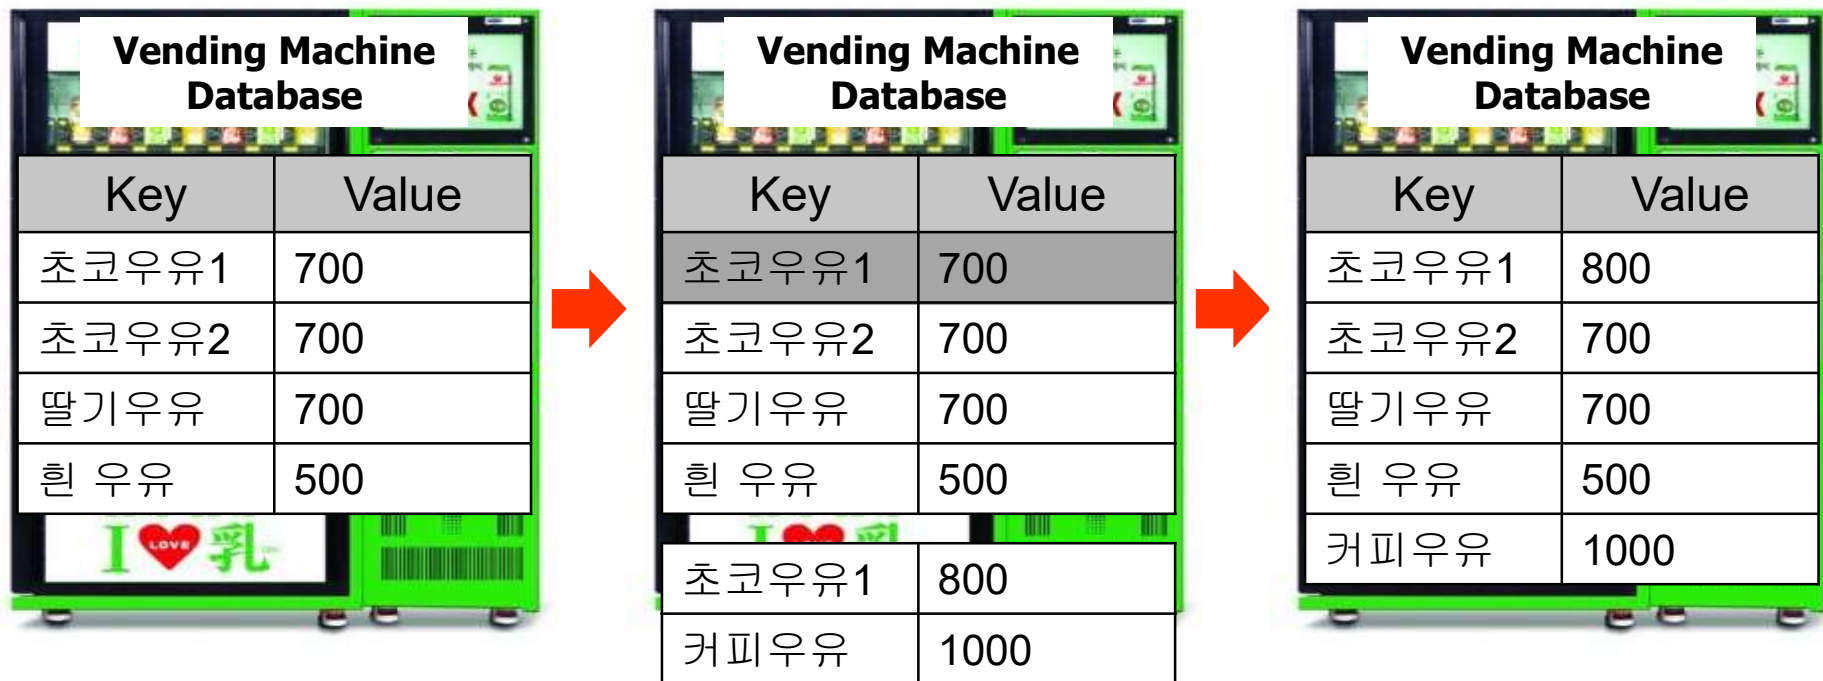
Candidate 1

Key	Value
700	초코, 딸기
500	흰 우유
...	...
...	...

Candidate 2

RocksDB Festival: Overall

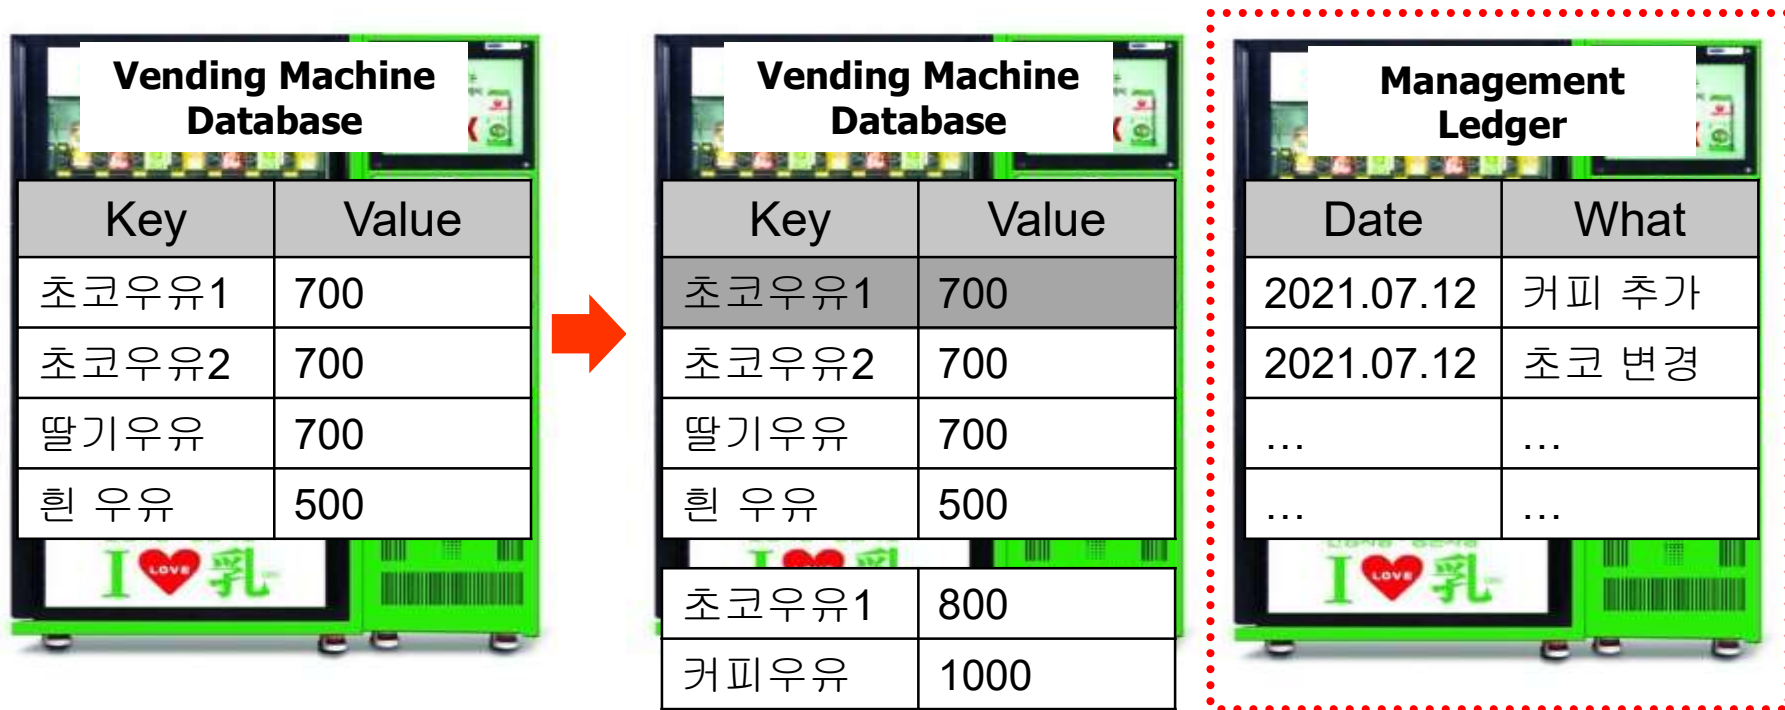
■ RocksDB Architecture: Vending Machine



**Compaction
(Merge Sort)**

RocksDB Festival: Overall

■ RocksDB Architecture: Vending Machine




WAL
(Write-Ahead-Log)

RocksDB Festival: MemTable

■ RocksDB Architecture: MemTable

- ✓ In-memory data structure
- ✓ MemTable full → Immutable memtable (Read-only)
- ✓ Type before being flushed to SSTable
- ✓ Immutable memtable flushed → delete memtable (free)



Input Data

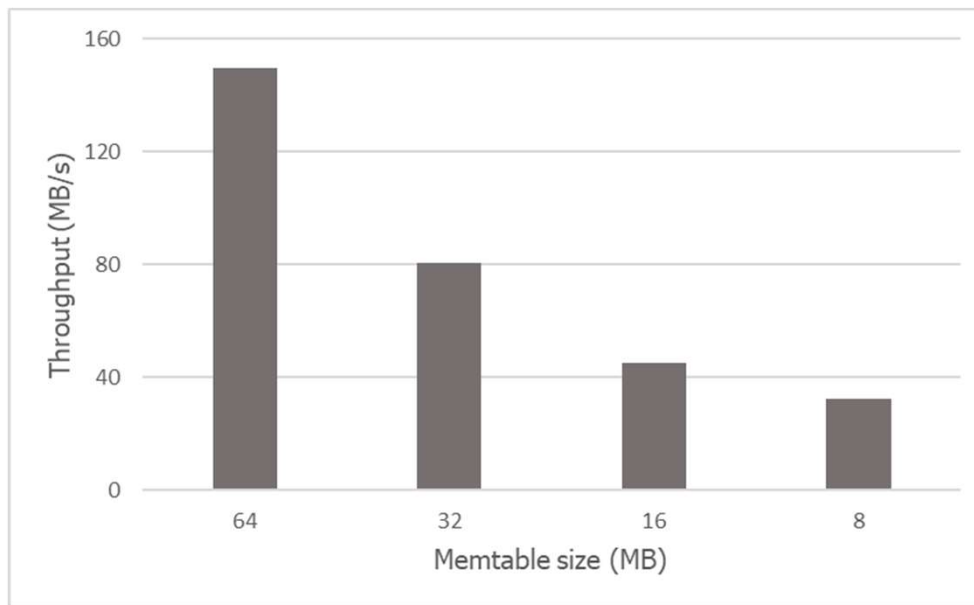


MemTable

RocksDB Festival: MemTable

■ RocksDB Architecture: MemTable

- ✓ Effect of MemTable Size: With RocksDB db_bench
- ✓ Why → Compaction Overhead




RocksDB Festival: SSTable

■ RocksDB Architecture: SSTable

- ✓ Sorted String Table
- ✓ Block, PlainTable, CuckooTable



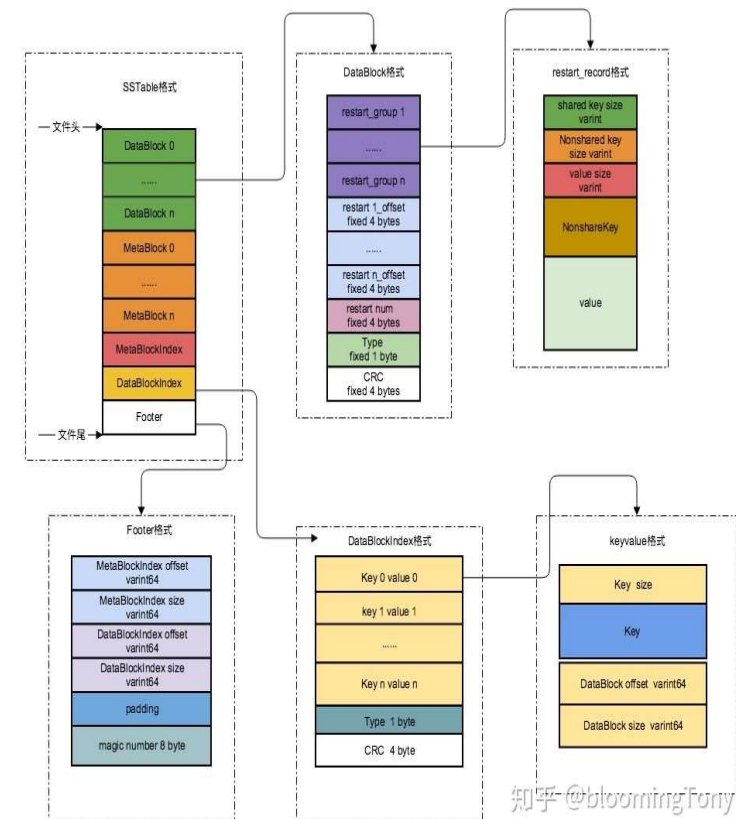
...



...



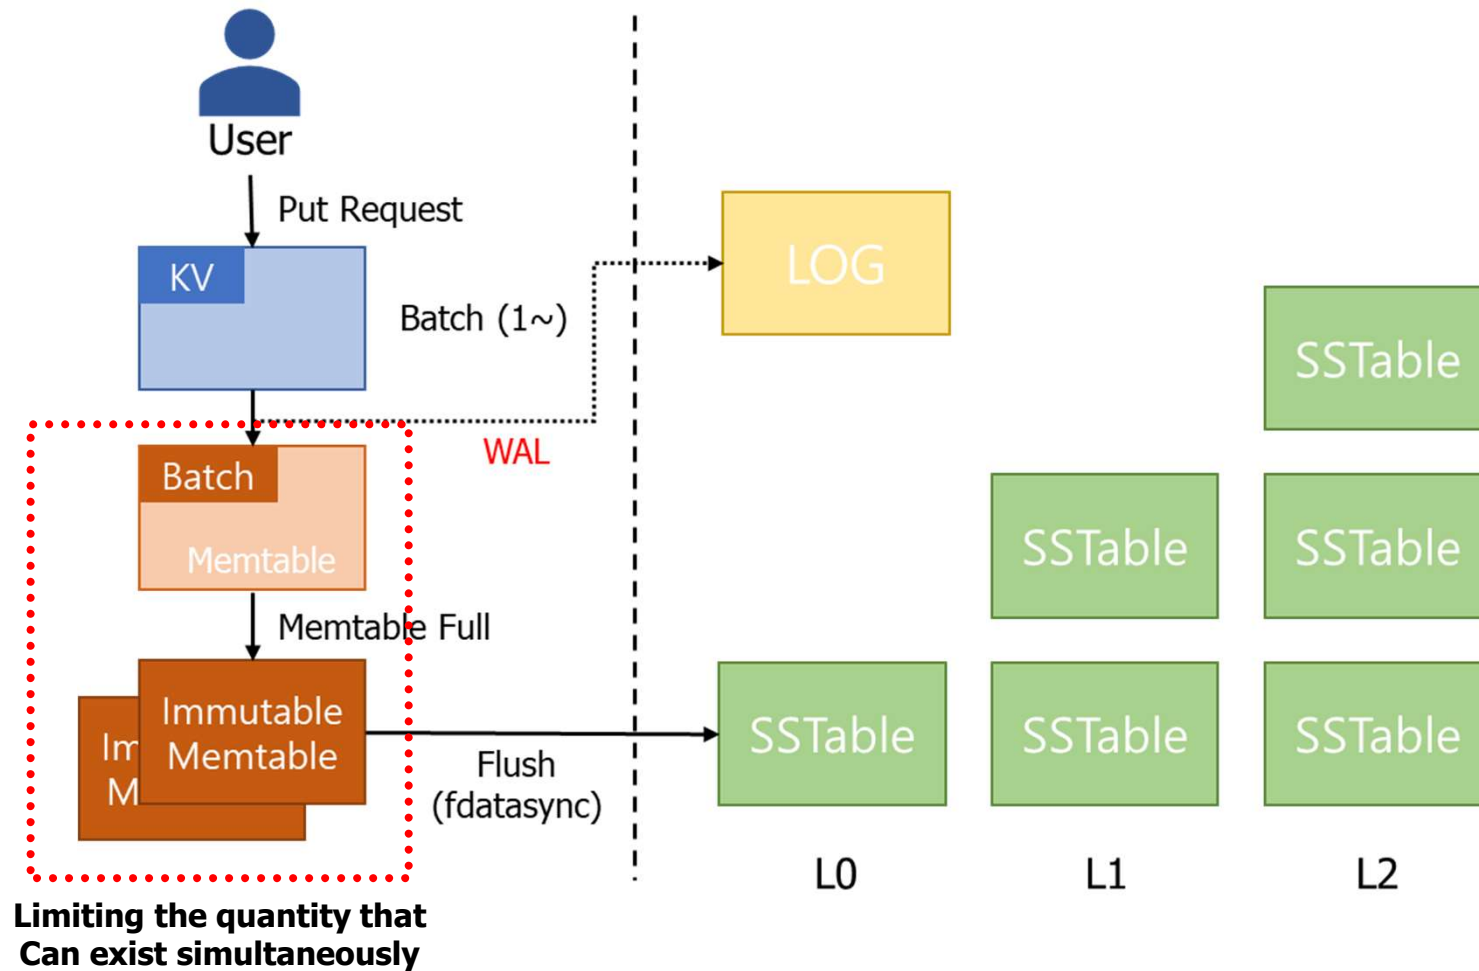
Filter block
Index block
Compression block
Range deletion block
Stats block



知乎@bloomingFony

RocksDB Festival: Write and WAL

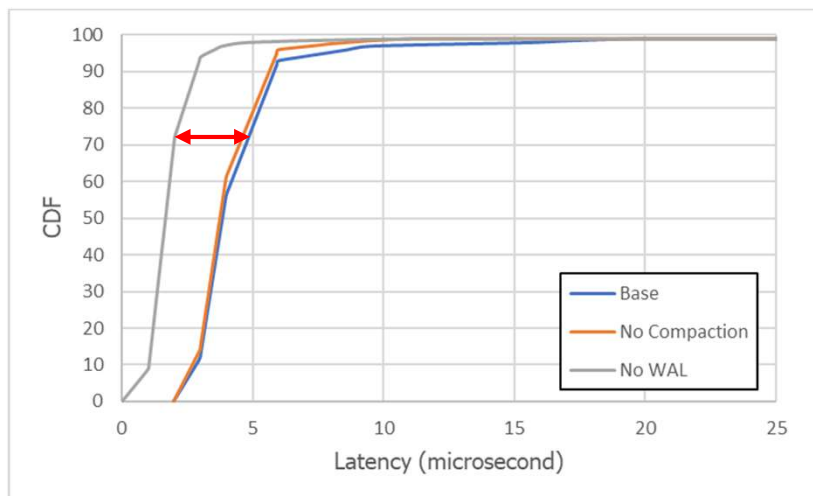
■ RocksDB Operation: Write and WAL



RocksDB Festival: Write and WAL

■ RocksDB Operation: Write and WAL


- ✓ 0000~~.log ➔ log file name
- ✓ 0000~~.sst ➔ SSTable file name

[illegible]


 **WAL overhead
(Write Latency)**

RocksDB Festival: Write and WAL

- RocksDB Operation: Write and WAL
 - ✓ If system failures can be predicted in advance
 - ✓ Control WAL overhead



**Do not proceed with
WAL operation**

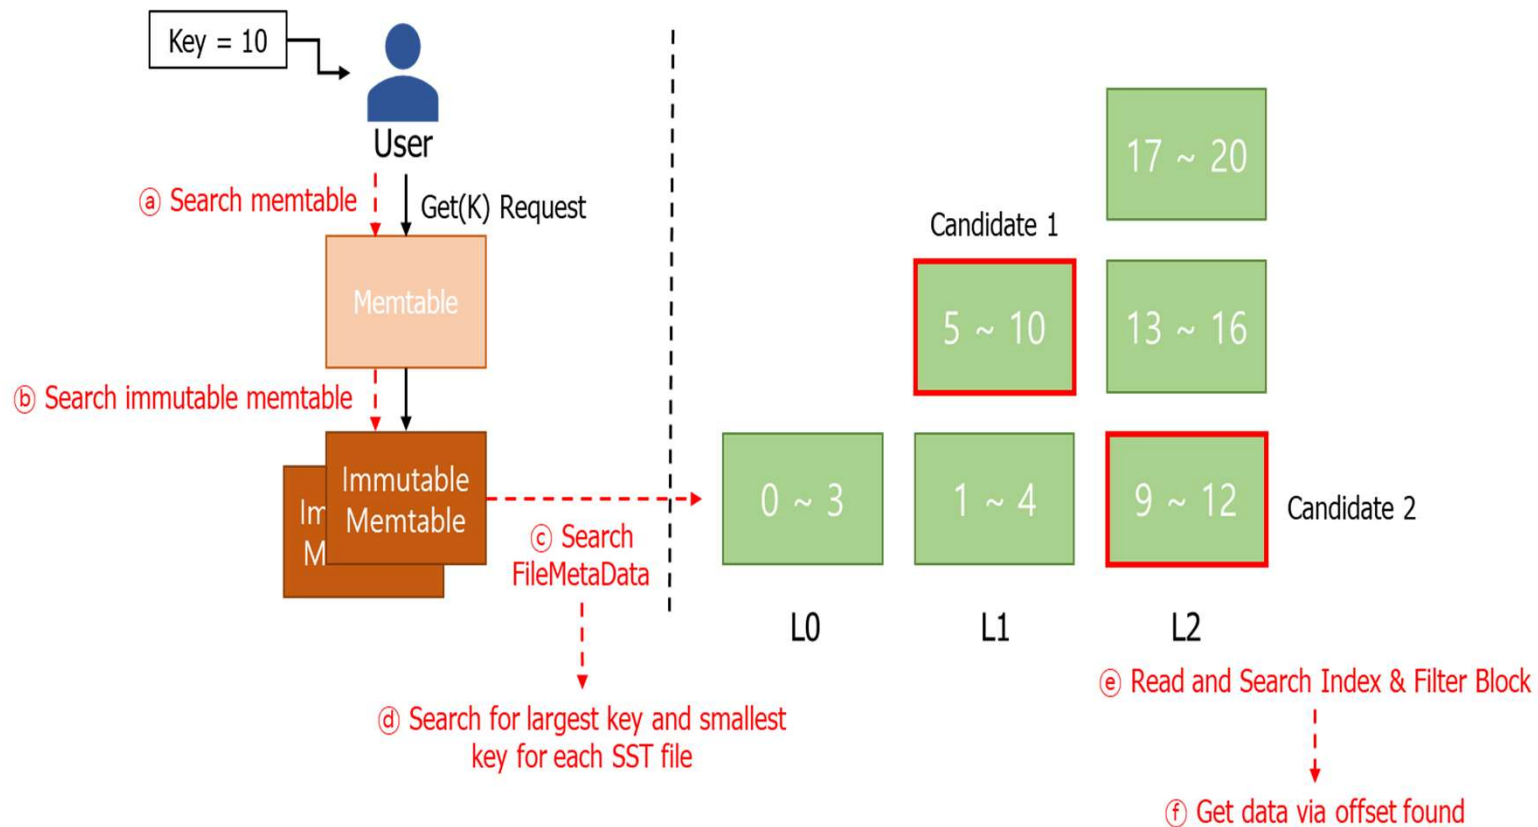


**Do proceed with
WAL operation**



RocksDB Festival: Read

■ RocksDB Operation: Read




RocksDB Festival: Read


■ RocksDB Operation: Read

- ✓ Index block, Bloom filter block


Case 1: K1



Case 2: K1

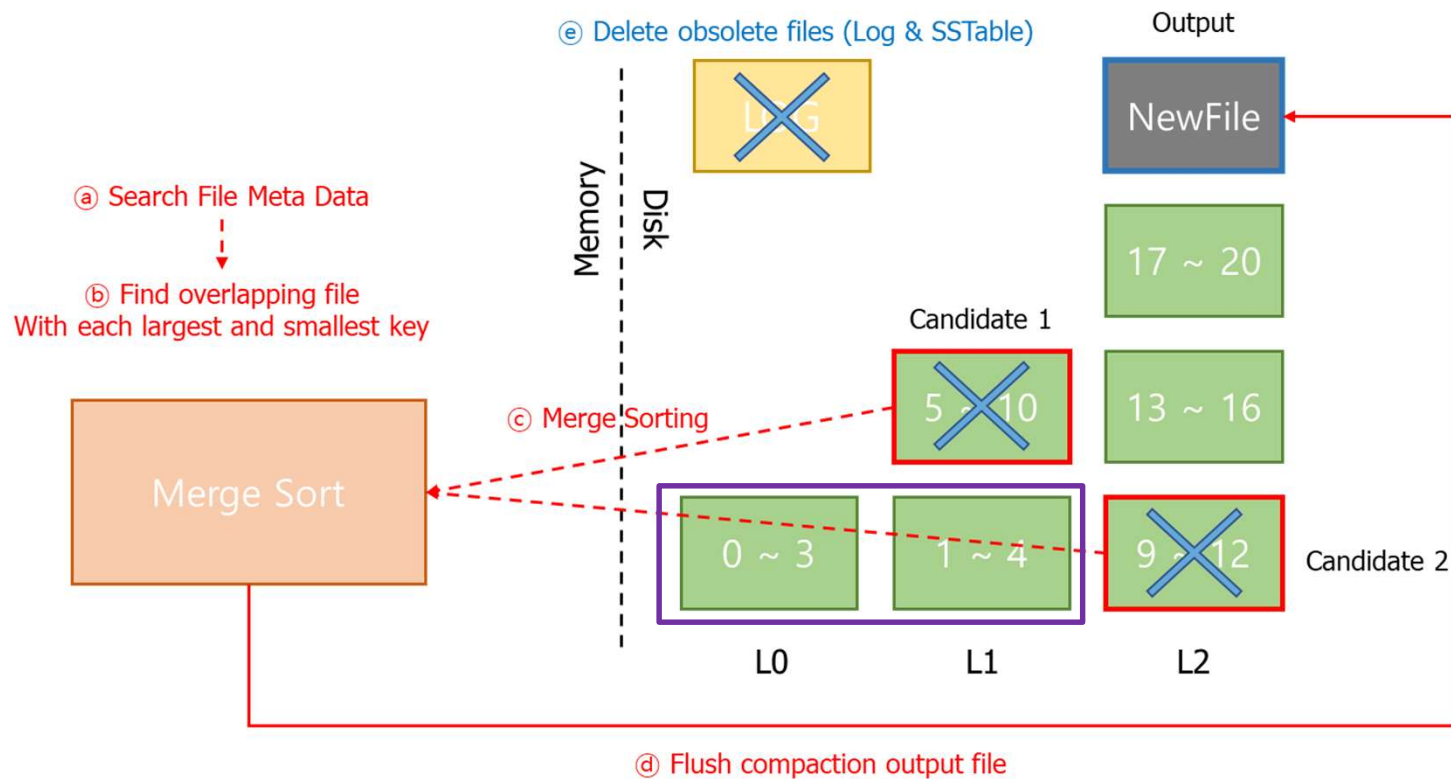


Case 3: K1



RocksDB Festival: Compaction

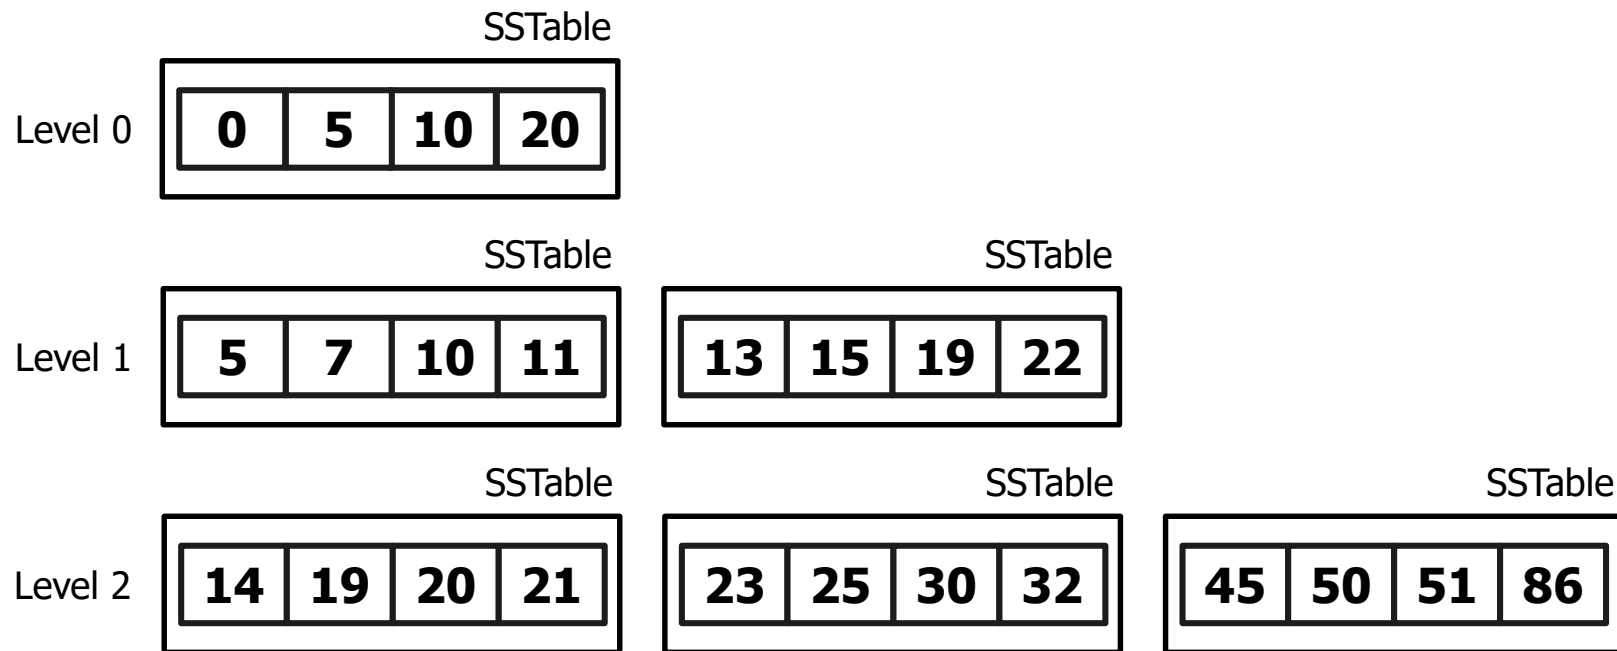
■ RocksDB Operation: Compaction



RocksDB Festival: Compaction

■ RocksDB Operation: Compaction

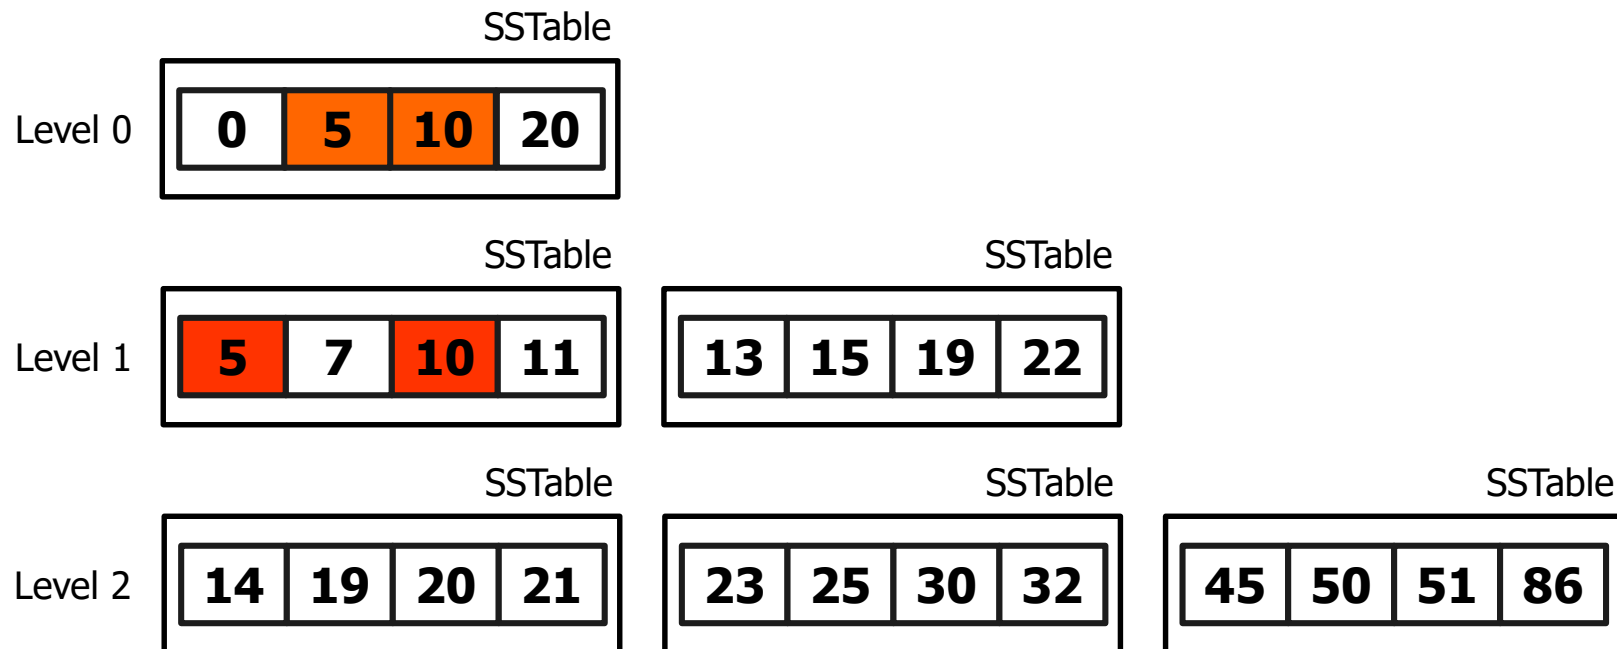
- ✓ Leveled Compaction, Universal Compaction, FIFO compaction
- ✓ Write Amplification, Space Amplification



RocksDB Festival: Compaction

■ RocksDB Operation: Compaction

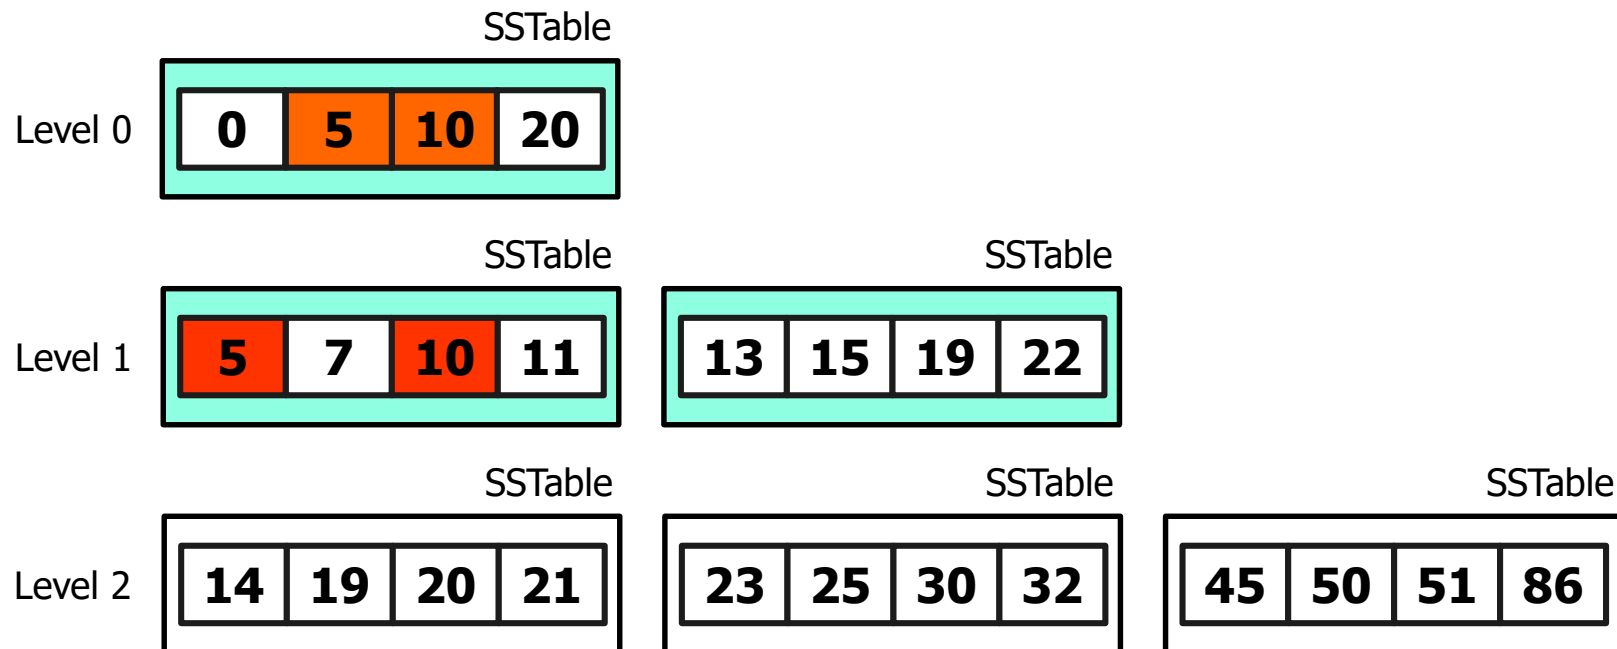
- ✓ Leveled Compaction, Universal Compaction, FIFO compaction
- ✓ Write Amplification, Space Amplification



RocksDB Festival: Compaction

■ RocksDB Operation: Compaction

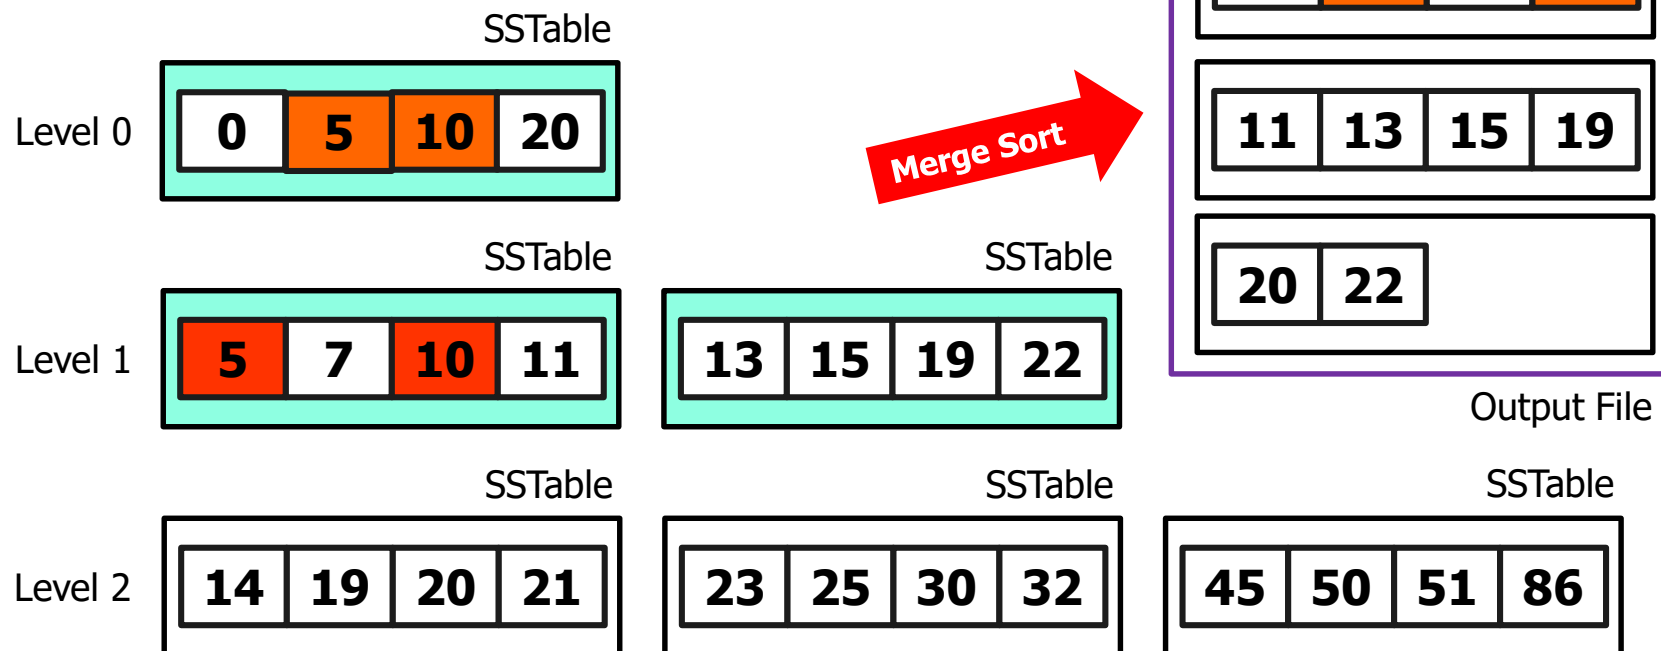
- ✓ Leveled Compaction, Universal Compaction, FIFO compaction
- ✓ Write Amplification, Space Amplification



RocksDB Festival: Compaction

■ RocksDB Operation: Compaction

- ✓ Leveled Compaction, Universal Compaction, FIFO compaction
- ✓ Write Amplification, Space Amplification

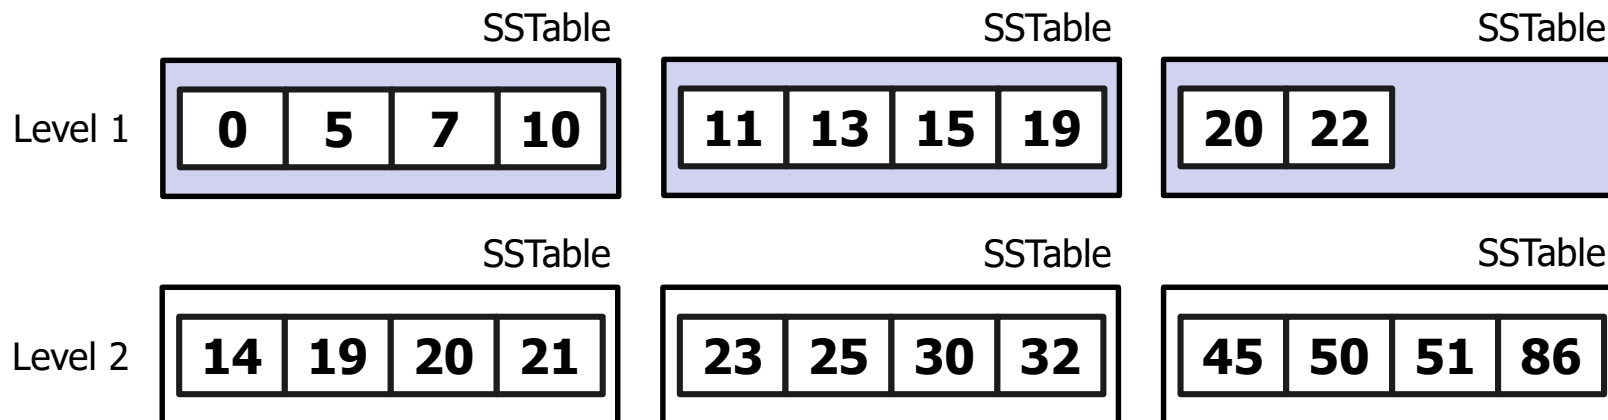


RocksDB Festival: Compaction

■ RocksDB Operation: Compaction

- ✓ Leveled Compaction, Universal Compaction, FIFO compaction
- ✓ Write Amplification, Space Amplification

Level 0

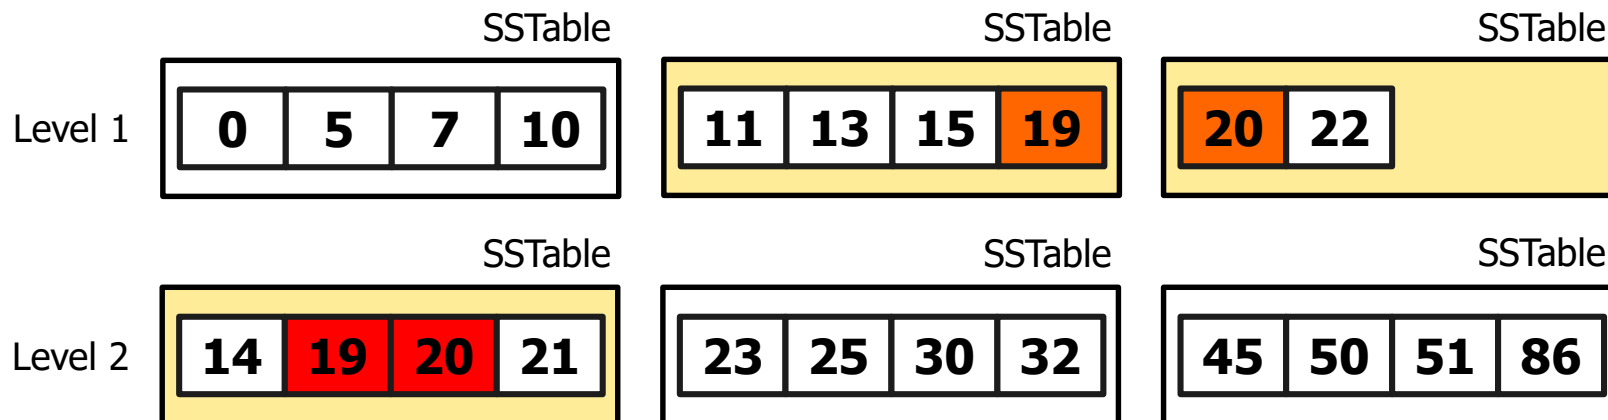


RocksDB Festival: Compaction

■ RocksDB Operation: Compaction

- ✓ Leveled Compaction, Universal Compaction, FIFO compaction
- ✓ Write Amplification, Space Amplification

Level 0




RocksDB Festival: Compaction

■ RocksDB Operation: Compaction

- ✓ Leveled Compaction, Universal Compaction, FIFO compaction
- ✓ Write Amplification, Space Amplification

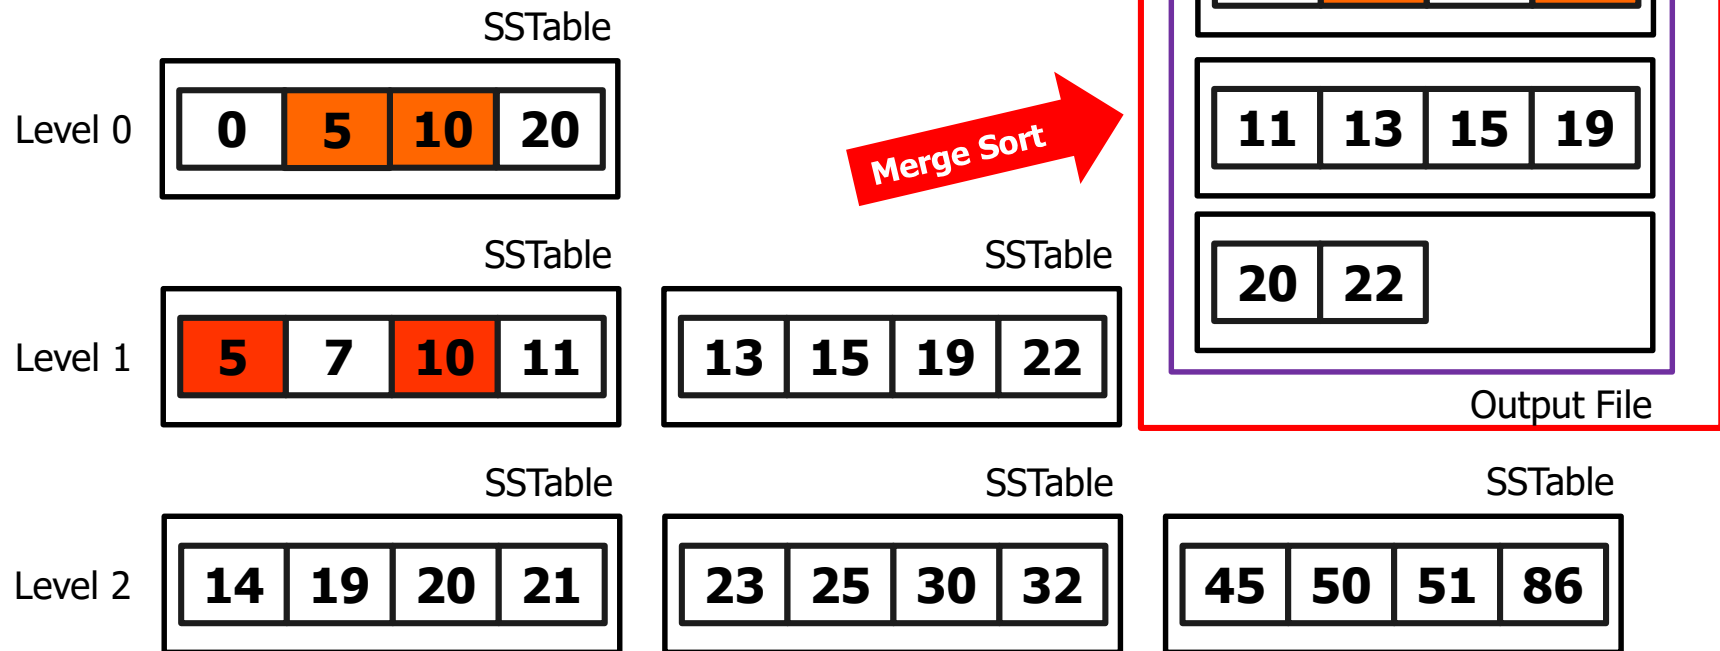
Level 0



RocksDB Festival: Compaction


■ RocksDB Operation: Compaction

✓ Write Amplification → $WAF = \frac{\text{Bytes written to Storage}}{\text{Bytes written to Database}}$



RocksDB Festival: Compaction

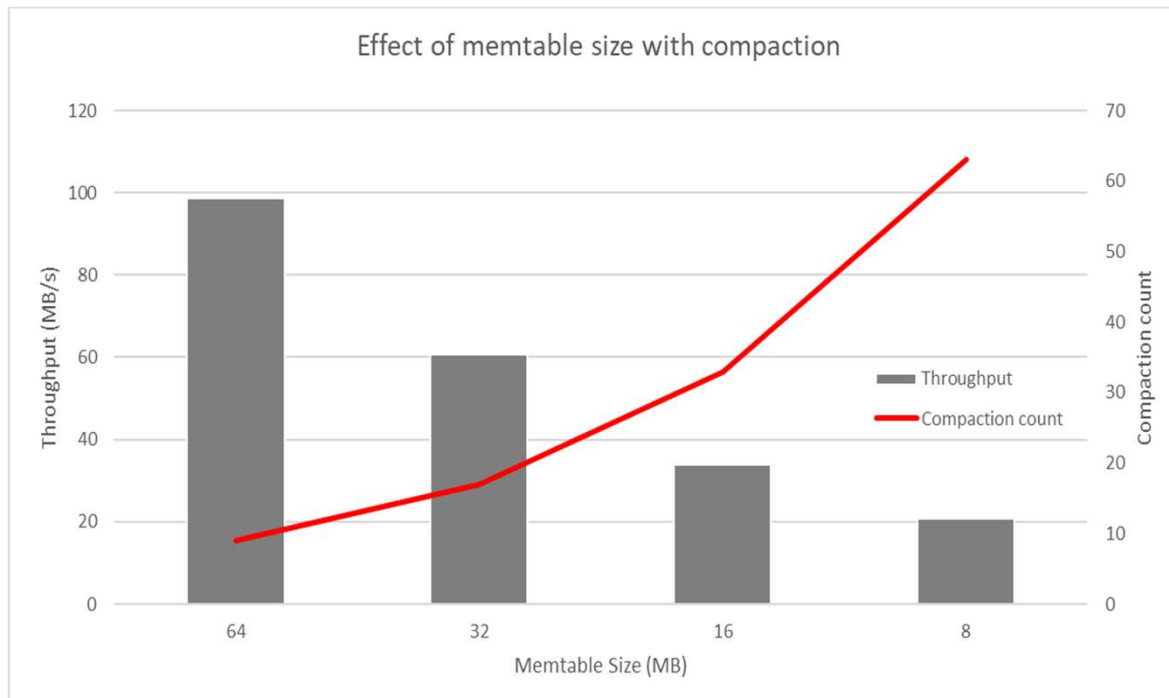
- RocksDB Operation: Compaction
 - ✓ Space Amplification



RocksDB Festival: Compaction

■ RocksDB Operation: Compaction

✓ Compaction Overhead



RocksDB Festival: Compaction

■ RocksDB Operation: Compaction

✓ Compaction Trigger



RocksDB Festival

■ Next Week

- ✓ 각 팀 발표 준비
 - 팀 구성
 - RocksDB db_bench 설치 및 실험 결과 (관심 분야)
 - 선택한 주제와 선택한 이유
 - 이번 스터디 동안 기대하는 결과

QnA

