

RocksDB Festival

RocksDB Env

Supported by IITP, StarLab.

July 5, 2021

Hojin Shin, Jongmoo Choi

choijm@dankook.ac.kr

<http://embedded.dankook.ac.kr/~choijm>

RocksDB Festival

■ Git clone RocksDB

- ✓ git clone <https://github.com/facebook/rocksdb.git>
- ✓ cd rocksdb

```
root@shin96:/home/shin96/RocksDB_Festival# git clone https://github.com/facebook/rocksdb.git
Cloning into 'rocksdb'...
remote: Enumerating objects: 99699, done.
remote: Counting objects: 100% (589/589), done.
remote: Compressing objects: 100% (296/296), done.
remote: Total 99699 (delta 308), reused 465 (delta 272), pack-reused 99110
Receiving objects: 100% (99699/99699), 168.23 MiB | 22.59 MiB/s, done.
Resolving deltas: 100% (75591/75591), done.
Checking connectivity... done.
root@shin96:/home/shin96/RocksDB_Festival# ls
rocksdb
root@shin96:/home/shin96/RocksDB_Festival# cd rocksdb/
root@shin96:/home/shin96/RocksDB_Festival/rocksdb# ls
appveyor.yml      db                HISTORY.md        memtable          test_util
AUTHORS           db_stress_tool   include           monitoring        third-party
buckifier         DEFAULT_OPTIONS_HISTORY.md  INSTALL.md       options          thirdparty.inc
build_tools       defs.bz1         issue_template.md plugin            tools
cache            docs             java             PLUGINS.md       trace_replay
cmake            DUMP_FORMAT.md  LANGUAGE-BINDINGS.md  port            USERS.md
CMakeLists.txt   env             LICENSE.Apache    README.md        util
CODE_OF_CONDUCT.md  examples        LICENSE.leveldb   ROCKSDB_LITE.md  utilities
CONTRIBUTING.md  file           logging          src.mk           Vagrantfile
COPYING          fuzz           Makefile         table            WINDOWS_PORT.md
coverage         hdfs          memory          TARGETS
```

RocksDB Festival

- Git clone RocksDB cont'
 - ✓ make db_bench
 - ✓ Then you can use RocksDB db_bench
 - ✓ It takes long time

```
root@shin96:/home/shin96/RocksDB_Festival/rocksdb# make db_bench
$DEBUG_LEVEL is 1
Makefile:176: Warning: Compiling in debug mode. Don't use the resulting binary in production
$DEBUG_LEVEL is 1
Makefile:176: Warning: Compiling in debug mode. Don't use the resulting binary in production
CC      tools/db_bench.o
CC      tools/db_bench_tool.o
CC      tools/simulated_hybrid_file_system.o
CC      test_util/testutil.o
CC      cache/cache.o
CC      cache/cache_entry_roles.o
CC      cache/clock_cache.o
CC      cache/lru_cache.o
CC      cache/sharded_cache.o
CC      db/arena_wrapped_db_iter.o
CC      db/blob/blob_fetcher.o
_
```

RocksDB Festival

■ Git clone RocksDB cont'

- ✓ ./db_bench ~~something options
- ✓ If you want set your device, you use [--db=(DEV_PATH)] option
- ✓ Example)

```
root@shin96:/home/shin96/RocksDB_Festival/rocksdb# ls
appveyor.yml      coverage          file              LICENSE.Apache   plugin            thirdparty.inc
AUTHORS            db               fuzz              LICENSE.leveldb  PLUGINS.md        tools
buckifier          db_bench         hdfs              logging           port              trace_replay
build_tools       db_stress_tool  HISTORY.md        make_config.mk   README.md         USERS.md
cache              DEFAULT_OPTIONS_HISTORY.md  include           Makefile         ROCKSDB_LITE.md  util
cmake              defs.bz2         INSTALL.md        memory            src.mk            utilities
CMakeLists.txt    docs            issue_template.md memtable          table             Vagrantfile
CODE_OF_CONDUCT.md DUMP_FORMAT.md  LANGUAGE-BINDINGS.md mnt              TARGETS           WINDOWS_PORT.md
CONTRIBUTING.md  env              librocksdb_debug.a monitoring         test_util
COPYING            examples         rocksdb_debug.a  options           thirdparty

root@shin96:/home/shin96/RocksDB_Festival/rocksdb# ./db_bench --db=../mnt --benchmarks=fillseq --num=1000000
Initializing RocksDB Options from the specified file
Initializing RocksDB Options from command-line flags
RocksDB: version 6.22
Date:      Fri Jul  2 22:02:50 2021
CPU:       4 * Intel(R) Core(TM) i5-6600 CPU @ 3.30GHz
CPUCache:  6144 KB
Keys:      16 bytes each (+ 0 bytes user-defined timestamp)
Values:    100 bytes each (50 bytes after compression)
Entries:   1000000
Prefix:    0 bytes
Keys per prefix: 0
RawSize:   110.6 MB (estimated)
FileSize:  62.9 MB (estimated)
Write rate: 0 bytes/second
Read rate: 0 ops/second
Compression: Snappy
Compression sampling rate: 0
Memtablerep: skip_list
Perf Level: 1
WARNING: Assertions are enabled; benchmarks unnecessarily slow
-----
Initializing RocksDB Options from the specified file
Initializing RocksDB Options from command-line flags
DB path: [../mnt]
fillseq    :          2.606 micros/op 383782 ops/sec;  42.5 MB/s
root@shin96:/home/shin96/RocksDB_Festival/rocksdb#
```

RocksDB Festival

■ Git clone RocksDB cont'

- ✓ You can find db_bench options in tools/db_bench_tool.cc
- ✓ Example)

```
DEFINE_string(column_family_distribution, "",
    "Comma-separated list of percentages, where the ith element "
    "indicates the probability of an op using the ith column family. "
    "The number of elements must be `num_hot_column_families` if "
    "specified; otherwise, it must be `num_column_families`. The "
    "sum of elements must be 100. E.g., if `num_column_families=4`, "
    "and `num_hot_column_families=0`, a valid list could be "
    "\"10,20,30,40\".");

DEFINE_int64(reads, -1, "Number of read operations to do. "
    "If negative, do FLAGS_num reads.");

DEFINE_int64(deletes, -1, "Number of delete operations to do. "
    "If negative, do FLAGS_num deletions.");

DEFINE_int32(bloom_locality, 0, "Control bloom filter probes locality");

DEFINE_int64(seed, 0, "Seed base for random number generators. "
    "When 0 it is deterministic.");

DEFINE_int32(threads, 1, "Number of concurrent threads to run.");

DEFINE_int32(duration, 0, "Time in seconds for the random-ops tests to run. "
    "When 0 then num & reads determine the test duration");

DEFINE_string(value_size_distribution_type, "fixed",
    "Value size distribution type: fixed, uniform, normal");

DEFINE_int32(value_size, 100, "Size of each value in fixed distribution");
static unsigned int value_size = 100;

DEFINE_int32(value_size_min, 100, "Min size of random value");

DEFINE_int32(value_size_max, 102400, "Max size of random value");

DEFINE_int32(seek_nexts, 0,
    "How many times to call Next() after Seek() in "
    "fillseekseq, seekrandom, seekrandomwhilewriting and "
    "seekrandomwhilemerging");

DEFINE_bool(reverse_iterator, false,
    "When true use Prev rather than Next for iterators that do "
    "Seek and then Next");

DEFINE_int64(max_scan_distance, 0,
    "Used to define iterate_upper_bound (or iterate_lower_bound "
    "if FLAGS_reverse_iterator is set to true) when value is nonzero");

DEFINE_bool(use_uint64_comparator, false, "use Uint64 user comparator");
```


RocksDB Festival

■ Git clone RocksDB cont'

- ✓ After modifying the source code, compilation must be performed again through make db_bench
- ✓ Example)

```
void Run() {  
    if (!SanityCheck()) {  
        ErrorExit();  
    }  
    Open(&open_options_);  
    PrintHeader();  
    fprintf(stdout, "Source Code Modifying!!!\n");  
    std::stringstream benchmark_stream(FLAGS_benchmarks);
```

```
root@shin96:/home/shin96/RocksDB_Festival/rocksdb# make db_bench  
$DEBUG_LEVEL is 1  
Makefile:176: Warning: Compiling in debug mode. Don't use the resulting binary in production  
$DEBUG_LEVEL is 1  
Makefile:176: Warning: Compiling in debug mode. Don't use the resulting binary in production  
CC      tools/db_bench_tool.o  
CCLD    db_bench  
root@shin96:/home/shin96/RocksDB_Festival/rocksdb#
```

```
root@shin96:/home/shin96/RocksDB_Festival/rocksdb# ./db_bench --db=  
Initializing RocksDB Options from the specified file  
Initializing RocksDB Options from command-line flags  
RocksDB: version 6.22  
Date: Fri Jul 2 22:06:32 2021  
CPU: 4 * Intel(R) Core(TM) i5-6600 CPU @ 3.30GHz  
CPU Cache: 6144 KB  
Keys: 16 bytes each (+ 0 bytes user-defined timestamp)  
Values: 100 bytes each (50 bytes after compression)  
Entries: 1000000  
Prefix: 0 bytes  
Keys per prefix: 0  
RawSize: 110.6 MB (estimated)  
FileSize: 62.9 MB (estimated)  
Write rate: 0 bytes/second  
Read rate: 0 ops/second  
Compression: Snappy  
Compression sampling rate: 0  
Memtable: skip_list  
Perf Level: 1  
WARNING: Assertions are enabled; benchmarks unnecessarily slow  
-----  
Source Code Modifying!!!  
Initializing RocksDB Options from the specified file  
Initializing RocksDB Options from command-line flags  
DB path: [./mnt]  
fillseq : 2.649 micros/op 377562 ops/sec; 41.8 MB/s  
root@shin96:/home/shin96/RocksDB_Festival/rocksdb#
```

RocksDB Festival

■ YCSB

- ✓ git clone https://github.com/brianfrankcooper/YCSB.git
- ✓ cd YCSB

```
root@shin96:/home/shin96/RocksDB_Festival# git clone https://github.com/brianfrankcooper/YCSB.git
Cloning into 'YCSB'...
remote: Enumerating objects: 20648, done.
remote: Total 20648 (delta 0), reused 0 (delta 0), pack-reused 20648
Receiving objects: 100% (20648/20648), 31.68 MiB | 17.69 MiB/s, done.
Resolving deltas: 100% (8016/8016), done.
Checking connectivity... done.
root@shin96:/home/shin96/RocksDB_Festival# ls
rocksdb  YCSB
root@shin96:/home/shin96/RocksDB_Festival# cd YCSB/
root@shin96:/home/shin96/RocksDB_Festival/YCSB# ls
accumulo1.9  cassandra  distribution  googledatastore  LICENSE.txt  pom.xml  s3  zookeeper
aerospike    checkstyle.xml  doc  griddb  maprdb  postgresql  scylla
arangodb      cloudspanner  dynamodb  hbase1  maprjsondb  rados  seaweedfs
asynchbase    CONTRIBUTING.md  elasticsearch  hbase2  memcached  README.md  solr7
azurecosmos    core  elasticsearch5  ignite  mongodb  redis  tablestore
azuresqlstorage  couchbase  foundationdb  infinispan  nosqlldb  rest  tarantool
bin            couchbase2  geode  jdbc  NOTICE.txt  riak  voltdb
binding-parent  crail  googlebigtable  kudu  orientdb  rocksdb  workloads
root@shin96:/home/shin96/RocksDB_Festival/YCSB#
```

RocksDB Festival

■ YCSB cont'

- ✓ mvn -pl sit.ycsb:rocksdb-binding -am clean package

```
root@shin96:/home/shin96/RocksDB_Festival/YCSB# mvn -pl site.ycsb:rocksdb-binding -am clean package
[INFO] Scanning for projects...
[INFO] -----
[INFO] Reactor Build Order:
[INFO]
[INFO] YCSB Root
[INFO] Core YCSB
[INFO] Per Datastore Binding descriptor
[INFO] YCSB Datastore Binding Parent
[INFO] RocksDB Java Binding
[INFO] -----
[INFO] Building YCSB Root 0.18.0-SNAPSHOT
[INFO] -----
[INFO]
[INFO] --- maven-clean-plugin:2.5:clean (default-clean) @ root ---
[INFO] Deleting /home/shin96/RocksDB_Festival/YCSB/target
[INFO]
[INFO] --- maven-enforcer-plugin:3.0.0-M1:enforce (enforce-maven) @ root ---
[INFO]
[INFO] --- maven-checkstyle-plugin:2.16:check (validate) @ root ---
[INFO]
[INFO] -----
[INFO] Building Core YCSB 0.18.0-SNAPSHOT
[INFO] -----
[INFO]
[INFO] --- maven-clean-plugin:2.5:clean (default-clean) @ core ---
[INFO] Deleting /home/shin96/RocksDB_Festival/YCSB/core/target
[INFO]
[INFO] --- maven-enforcer-plugin:3.0.0-M1:enforce (enforce-maven) @ core ---
[INFO]
[INFO] --- maven-checkstyle-plugin:2.16:check (validate) @ core ---
[INFO]
[INFO] --- maven-resources-plugin:2.6:resources (default-resources) @ core ---
[INFO] Using 'UTF-8' encoding to copy filtered resources.
[INFO] Copying 1 resource
[INFO]
[INFO] --- maven-compiler-plugin:3.7.0:compile (default-compile) @ core ---
[INFO] Changes detected - recompiling the module!
[INFO] Compiling 63 source files to /home/shin96/RocksDB_Festival/YCSB/core/target/classes
[INFO]
[INFO] --- maven-resources-plugin:2.6:testResources (default-testResources) @ core ---
[INFO] Using 'UTF-8' encoding to copy filtered resources.
[INFO] skip non existing resourceDirectory /home/shin96/RocksDB_Festival/YCSB/core/src/test/resources
[INFO]
```


RocksDB Festival

■ YCSB cont'

✓ load workload

- `./bin/ycsb load rocksdb -s -P workloads/workloada -p rocksdb.dir=[PATH] -p recordcount=[num]`

✓ run workload

- `./bin/ycsb run rocksdb -s -P workloads/workloada -p rocksdb.dir=[PATH] -p operationcount=[num]`

```
Command line: -db site.ycsb.db.rocksdb.RocksDBClient -s -P workloads/workloada -p rocksdb.dir=/tmp/rocksdb-data-ycsb -load
YCSB Client 0.18.0-SNAPSHOT
```

```
Loading workload...
Starting test.
SLF4J: Failed to load class "org.slf4j.impl.StaticLoggerBinder".
SLF4J: Defaulting to no-operation (NOP) logger implementation
SLF4J: See http://www.slf4j.org/codes.html#StaticLoggerBinder for further details.
2021-07-02 22:00:25:218 0 sec: 0 operations; est completion in 0 second
DBWrapper: report latency for each error is false and specific error codes to track for latency are: []
2021-07-02 22:00:25:457 0 sec: 1000 operations; 3958.25 current ops/sec; [CLEANUP: Count=1, Max=1039, Min=1039, Avg=1039
, 90=1039, 99=1039, 99.9=1039, 99.99=1039] [INSERT: Count=1000, Max=52767, Min=11, Avg=78.21, 90=37, 99=162, 99.9=526, 9
9.99=52767]
[OVERALL], RunTime(ms), 253
[OVERALL], Throughput(ops/sec), 3952.5691699604745
[TOTAL_GCS_PS_Scavenge], Count, 0
[TOTAL_GC_TIME_PS_Scavenge], Time(ms), 0
[TOTAL_GC_TIME_PS_Scavenge], Time(%), 0.0
[TOTAL_GCS_PS_MarkSweep], Count, 0
[TOTAL_GC_TIME_PS_MarkSweep], Time(ms), 0
[TOTAL_GC_TIME_PS_MarkSweep], Time(%), 0.0
[TOTAL_GCS], Count, 0
[TOTAL_GC_TIME], Time(ms), 0
[TOTAL_GC_TIME_%], Time(%), 0.0
[CLEANUP], Operations, 1
[CLEANUP], AverageLatency(us), 1039.0
[CLEANUP], MinLatency(us), 1039
[CLEANUP], MaxLatency(us), 1039
[CLEANUP], 95thPercentileLatency(us), 1039
[CLEANUP], 99thPercentileLatency(us), 1039
[INSERT], Operations, 1000
[INSERT], AverageLatency(us), 78.208
[INSERT], MinLatency(us), 11
[INSERT], MaxLatency(us), 52767
[INSERT], 95thPercentileLatency(us), 62
[INSERT], 99thPercentileLatency(us), 162
[INSERT], Return=OK, 1000
root@shin96:/home/shin96/RocksDB_Festival/YCSB#
```

```
Command line: -db site.ycsb.db.rocksdb.RocksDBClient -s -P workloads/workloada -p rocksdb.dir=/tmp/rocksdb-data-ycsb -run
YCSB Client 0.18.0-SNAPSHOT
```

```
Loading workload...
Starting test.
SLF4J: Failed to load class "org.slf4j.impl.StaticLoggerBinder".
SLF4J: Defaulting to no-operation (NOP) logger implementation
SLF4J: See http://www.slf4j.org/codes.html#StaticLoggerBinder for further details.
2021-07-02 22:00:56:630 0 sec: 0 operations; est completion in 0 second
DBWrapper: report latency for each error is false and specific error codes to track for latency are: []
2021-07-02 22:00:56:845 0 sec: 1000 operations; 4366.81 current ops/sec; [READ: Count=521, Max=184, Min=3, Avg=13.95, 90
=28, 99=74, 99.9=140, 99.99=184] [CLEANUP: Count=1, Max=1123, Min=1123, Avg=1123, 90=1123, 99=1123, 99.9=1123, 99.99=112
3] [UPDATE: Count=479, Max=4065, Min=19, Avg=55.03, 90=86, 99=233, 99.9=4065, 99.99=4065]
[OVERALL], RunTime(ms), 229
[OVERALL], Throughput(ops/sec), 4366.812227074236
[TOTAL_GCS_PS_Scavenge], Count, 0
[TOTAL_GC_TIME_PS_Scavenge], Time(ms), 0
[TOTAL_GC_TIME_PS_Scavenge], Time(%), 0.0
[TOTAL_GCS_PS_MarkSweep], Count, 0
[TOTAL_GC_TIME_PS_MarkSweep], Time(ms), 0
[TOTAL_GC_TIME_PS_MarkSweep], Time(%), 0.0
[TOTAL_GCS], Count, 0
[TOTAL_GC_TIME], Time(ms), 0
[TOTAL_GC_TIME_%], Time(%), 0.0
[READ], Operations, 521
[READ], AverageLatency(us), 13.953934740882918
[READ], MinLatency(us), 3
[READ], MaxLatency(us), 184
[READ], 95thPercentileLatency(us), 41
[READ], 99thPercentileLatency(us), 74
[READ], Return=OK, 521
[CLEANUP], Operations, 1
[CLEANUP], AverageLatency(us), 1123.0
[CLEANUP], MinLatency(us), 1123
[CLEANUP], MaxLatency(us), 1123
[CLEANUP], 95thPercentileLatency(us), 1123
[CLEANUP], 99thPercentileLatency(us), 1123
[UPDATE], Operations, 479
[UPDATE], AverageLatency(us), 55.02713987473904
[UPDATE], MinLatency(us), 19
[UPDATE], MaxLatency(us), 4065
[UPDATE], 95thPercentileLatency(us), 153
[UPDATE], 99thPercentileLatency(us), 233
[UPDATE], Return=OK, 479
root@shin96:/home/shin96/RocksDB_Festival/YCSB#
```

Load workload

Run workload

RocksDB Festival

■ YCSB

- ✓ cd workloads
- ✓ You can control workload files

```
root@shin96:/home/shin96/RocksDB_Festival/YCSB/workloads# ls
tsworkloada tsworkload_template workloada workloadb workloadc workloadd worlkoade workloadf workload template
```

Discussion

