



# RocksDB Festival

## What is the RocksDB

Supported by IITP, StarLab.

July 5, 2021

Hojin Shin, Jongmoo Choi

[choijm@dankook.ac.kr](mailto:choijm@dankook.ac.kr)

<http://embedded.dankook.ac.kr/~choijm>

# RocksDB Festival: Contents

## ■ Contents

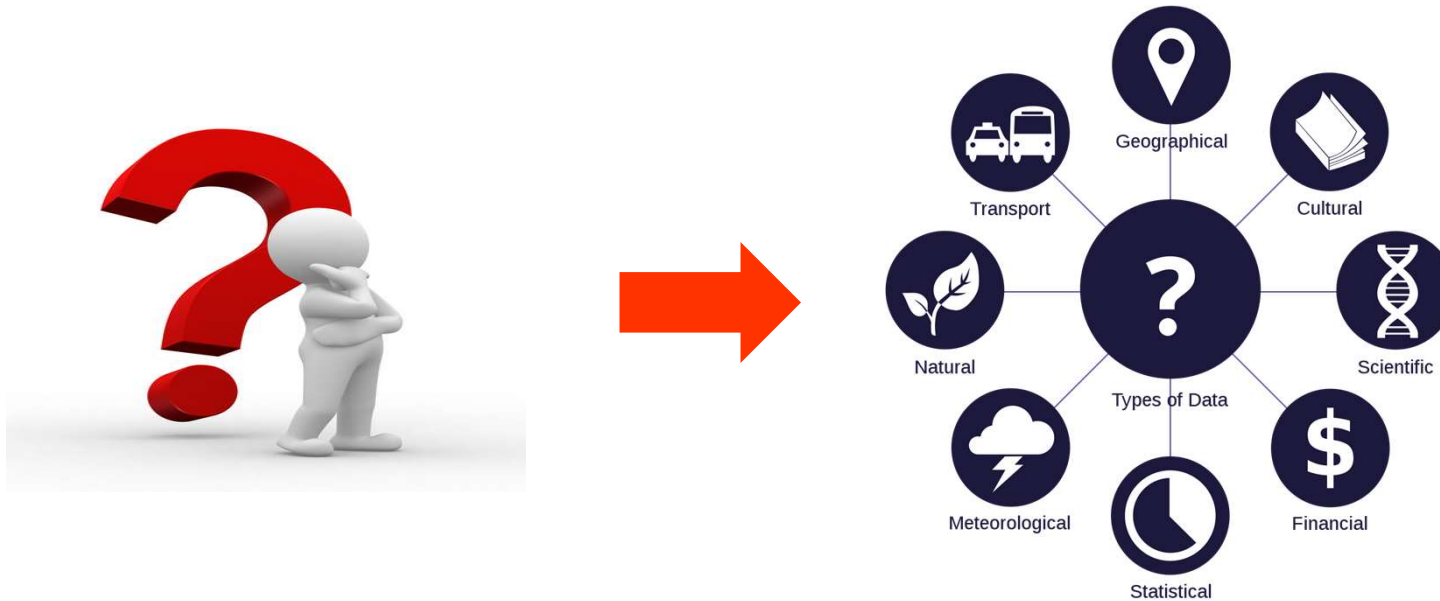
- ✓ What is data
- ✓ How to manage data
- ✓ This is Key-Value Store
- ✓ The basics of key value: LevelDB
- ✓ Our Goal: RocksDB



# RocksDB Festival: What is the Data

## ■ Data


- ✓ 1) Units of information, often numeric, that are collected through observation.
- ✓ 2) Fact on which a theory is based
- ✓ 3) Data in the form of letters, numbers, sounds, pictures that a computer can process




# RocksDB Festival: What is the Data

## ■ Information

- ✓ Information is obtained by processing data
- ✓ A form in which data is processed according to its meaning and purpose for specific decision-making




 <p>吃个汉堡吧</p>	 <p>汉堡包 18</p>	 <p>汉堡包 18</p>	 <p>汉堡包 18</p>	 <p>汉堡包 18</p>
 <p>汉堡包 18</p>	 <p>汉堡包 18</p>	 <p>冰淇淋 15</p>	 <p>冰淇淋 15</p>	 <p>冰淇淋 15</p>
 <p>冰淇淋 15</p>	 <p>冰淇淋 15</p>	 <p>冰淇淋 15</p>	 <p>冰淇淋 15</p>	 <p>冰淇淋 15</p>
 <p>冰淇淋 15</p>	 <p>冰淇淋 15</p>	 <p>冰淇淋 15</p>	 <p>冰淇淋 15</p>	 <p>冰淇淋 15</p>
 <p>冰淇淋 15</p>	 <p>冰淇淋 15</p>	 <p>冰淇淋 15</p>	 <p>冰淇淋 15</p>	 <p>冰淇淋 15</p>
 <p>冰淇淋 15</p>	 <p>冰淇淋 15</p>	 <p>冰淇淋 15</p>	 <p>冰淇淋 15</p>	 <p>冰淇淋 15</p>
 <p>冰淇淋 15</p>	 <p>冰淇淋 15</p>	 <p>冰淇淋 15</p>	 <p>冰淇淋 15</p>	 <p>冰淇淋 15</p>



# RocksDB Festival: What is the Data

## ■ Big Data

- ✓ A large amount of structured data that exceed existing DB management tools
- ✓ Set of unstructured data that is not in the form of data
- ✓ Features = 5V




# RocksDB Festival: What is the Data


## ■ Big Data Prospects



국내 빅데이터 및 분석 시장 전망 2019-2023년 [단위:십억]



Source: IDC, 2019



2011~2026년간 빅데이터 시장 규모 전망 (단위: 10억 달러)

출처: statista

### 대규모 사업체 신기술 이용 추이

250인 이상 사업체 기준



자료/ 과학기술정보통신부



김지영 인턴 / 20200121

트위터 @yonhap\_graphics 페이스북 tuneey.kr/LeYN1

# RocksDB Festival : What is the Data

---

## ■ Kind of Data

### ✓ Structured Data

- Data organized and processed into a form suitable for immediate statistical analysis
- Data stored in **fixed** fields

### ✓ Unstructured Data


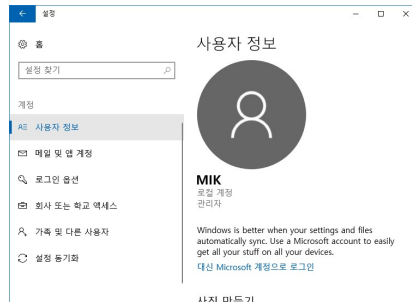
- One piece of data, **not a set of data**, is objectified as collected data
- Difficult to understand the meaning of a value because there is no set rule

### ✓ Semi-structured Data

- File type, metadata (schema of structured data inside data)

# RocksDB Festival: How to manage Data

## ■ Various Data



{JSON}



### 소비자 집단별 소비 형태 특성

	전통따라가치족	디지털헬스족	정보무장자기만족	바쁜미레족	유행선도족
계획적 VS. 충동적	계획적	충동적	충동적	계획적	계획적
광고 광고 VS. 정보 추구	정보 추구	광고 광고	정보 추구	정보 추구	광고 광고
매장 VS. 온라인	매장	온라인	온라인	온라인	매장
현금 VS. 카드	현금	현금	카드	카드	현금
가족 중심 VS. 개인 취향	가족 중심	개인 취향	개인 취향	가족 중심	개인 취향
자주 구매 VS. 몰아서 구매	자주 구매	자주 구매	몰아서 구매	몰아서 구매	몰아서 구매
먼저 구매 VS. 따라 구매	따라 구매	따라 구매	따라 구매	먼저 구매	먼저 구매

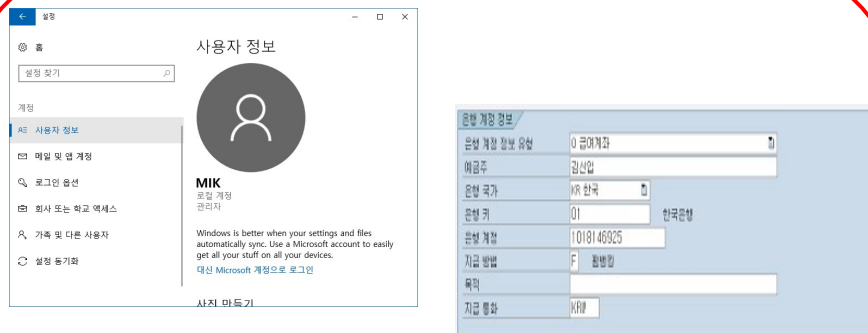
\* 집단별 특성에 따라 소비 형태에 대한 장점을 가진 특성을 기술함.





# RocksDB Festival: How to manage Data

## ■ Various Data



소비자 집단별 소비 형태 특성

전통대리점	디지털점포	정보통신망	비밀리에	유통선도
가정용 VS. 상업용	가정용	상업용	상업용	가정용
광고 광고 VS. 정보 추구	정보 추구	광고 광고	정보 추구	광고 광고
매장 VS. 온라인	매장	온라인	온라인	매장
현금 VS. 카드	현금	카드	카드	현금
가족 중심 VS. 개인 취향	가족 중심	개인 취향	개인 취향	가족 중심
자주 구매 VS. 몰에서 구매	자주 구매	몰에서 구매	몰에서 구매	몰에서 구매
편지 구매 VS. 택배 구매	편지 구매	택배 구매	택배 구매	편지 구매

\* 집단별 소비 형태에 따라 소비 형태를 구분하여 분류함

## Structured Data



## Semi-structured Data

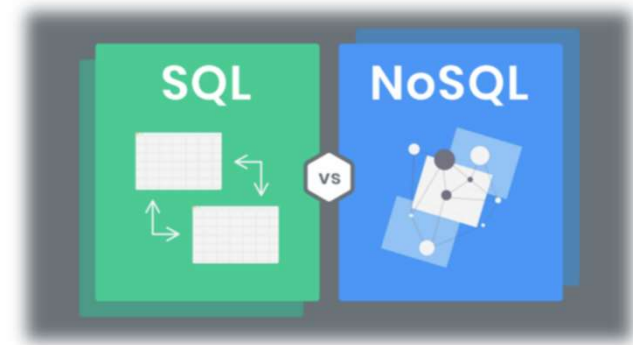


## Unstructured Data

# RocksDB Festival: How to manage Data

## ■ Data management

- ✓ SQL (Structured Query Language)
  - Interact with a particular type of database
  - Can store, modify, delete and retrieve data from RDBMS
  - Features
    - Strict schema
    - Relation
- ✓ NoSQL (Not only SQL)
  - Adjust the stored data at any time and add new "fields"
  - **Key-value**, document, wide-column, graph
  - Features
    - No schema
    - No Relation



# RocksDB Festival: How to manage Data

---

## ■ Data management cont'

### ✓ SQL

- RDBMS(Relational DataBase Management System)
- Easy to perform transactions by minimizing data redundancy through normalization
- Data integrity – Accuracy, consistence
- MSSQL, MySQL, Oracle

### ✓ NoSQL

- Does not define relationships between data
- Store large amounts of data
- LevelDB, RocksDB, Cassandra, MongoDB, Memcached



SQL

NoSQL

# RocksDB Festival: This is Key-Value Store

## ■ Key-Value Store

- ✓ A type of non-relational database that uses **key-value methods** to store data
- ✓ Key is a unique identifier and cannot overlapping value
- ✓ Value can be **anything** – integer, string, JSON, image ...
- ✓ Hash function is used to process the key

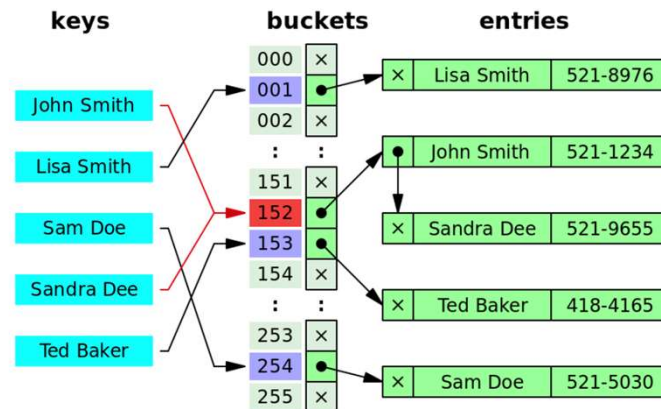
Key	Value
K1	AAA,BBB,CCC
K2	AAA,BBB
K3	AAA,DDD
K4	AAA,2,01/01/2015
K5	3,ZZZ,5623

### Key-Value Overview

# RocksDB Festival: This is Key-Value Store

## ■ Hash-Table

- ✓ Data structure that stores value in key
- ✓ When searching for data about a key, if you execute **hash function** only once → store and delete data is fast
- ✓ Mapping → Called Hashing
- ✓ Data access: insert, delete, retrieve  $O(1)$
- ✓ Problem
  - Hash collision → Can solve using “Chaining”




**Hash Table Overview**

# RocksDB Festival: This is Key-Value Store

## ■ Skiplist

- ✓ A data structure that enables fast search, insert, and delete with an algorithm applied to a sorted linked list
- ✓ Useful in multithreaded system architectures
- ✓ Complexity :  $O(\log n)$




**Skiplist Overview**

# RocksDB Festival: This is Key-Value Store

## ■ LSM(Log Structured Merge)-Tree

- ✓ Patrick O'Neil, The Log-Structured Merge Tree, 1996
- ✓ Has 0~L levels, L0 is in memory others are in disk(storage)
- ✓ Buffer located in L0 stores data and when buffer full, **flushed** to lower level one by one
- ✓ Not In-place-update, using Append
- ✓ Write
  - Memory buffer construct skiplist, maintain input data's order
- ✓ Read
  - When reading, all files are checked
  - Memory → Immutable memory → Disk

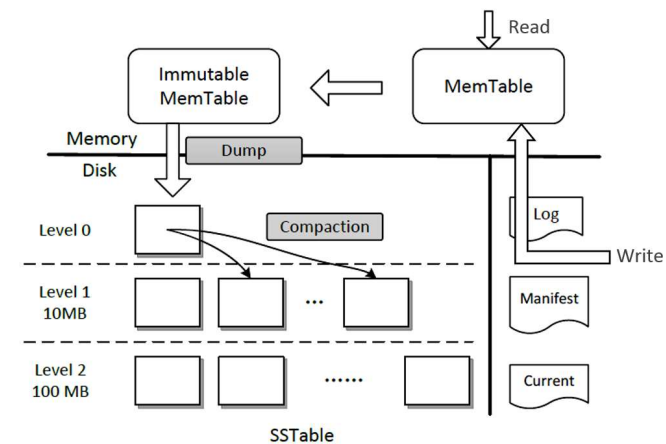


**LSM Tree Overview**

# RocksDB Festival: LevelDB

## ■ LevelDB

- ✓ Google's opensource project
- ✓ Developed in the programming language C++
- ✓ Data is stored after sorting by key
- ✓ Operation : Put(K, V), Get(K), Delete(K)
- ✓ Multiple operations can be created and processed in one batch
- ✓ Limitation
  - Single processing: only one process can access DB
  - Not support SQL query



## LevelDB Overview



# RocksDB Festival: RocksDB

---

## ■ Overview

- ✓ RocksDB is a storage engine for server workloads (Developed by Facebook)
- ✓ Data stored by key and value
- ✓ Flexibility: Support various production env. (Memory, Flash, HDD)
- ✓ Operation: Get(K), NewIterator(), Put(K,V), Delete(K) ...
- ✓ **Memtable**: In-memory data structure
- ✓ **Log**: Sequential write into storage
- ✓ **SSTable**: Sorted data (L1 >) in storage
- ✓ WAL(Write-Ahead-Log): Before flush memtable, do flush log

# RocksDB Festival: RocksDB

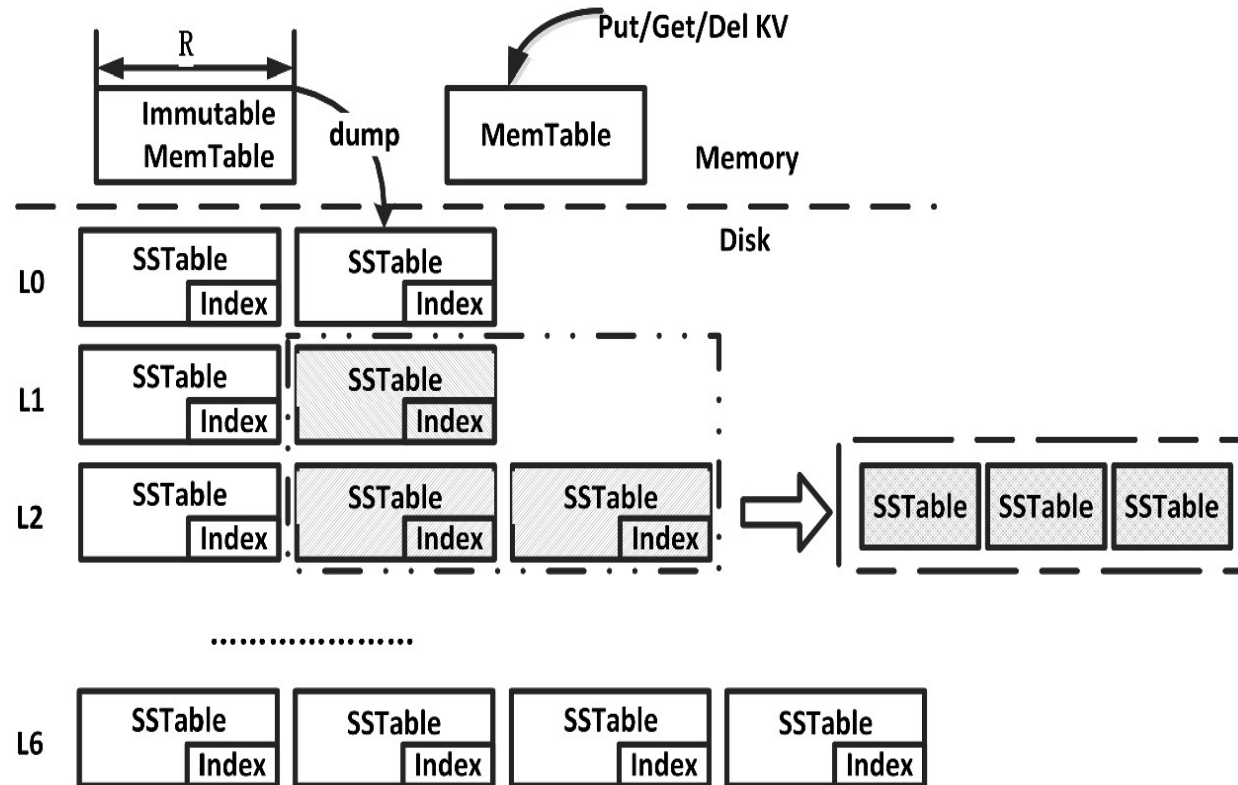
---

## ■ Overview cont'

- ✓ Column Family
  - Supports partitioning database into multiple column families
- ✓ Update
  - Put API inserts a single key-value into the database
- ✓ **Get**, Iterator
  - Get a single key-value from the database using the Get API
  - Iterator API allows applications to perform range scans on the database
- ✓ **Compaction**
  - Removes multiple copies of the same key that may occur when overwriting an existing key
  - Write throughput depends on compactions speed

# RocksDB Festival: RocksDB

## ■ RocksDB Architecture

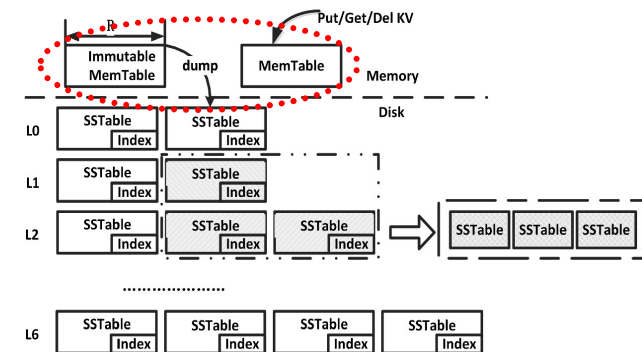


**RocksDB Architecture**

# RocksDB Festival: RocksDB

## ■ Main Terminology for RocksDB: MemTable

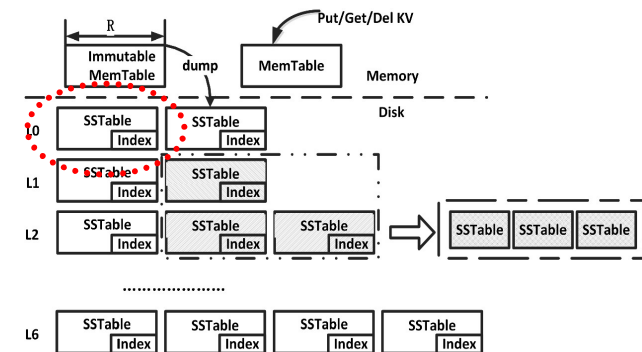
- ✓ Memtable is an **in-memory data-structure** that holds data before being flushed to the SST file
- ✓ When memtable full, it becomes immutable memtable
- ✓ The **background thread flushes** the contents of the memtable to the SST file
- ✓ Memtable has a skiplist structure
- ✓ Insert and check first when performing Put() and Get()
- ✓ Skiplist, HashSkiplist, HashLinklist, Vector



# RocksDB Festival: RocksDB

## ■ Main Terminology for RocksDB: SSTable

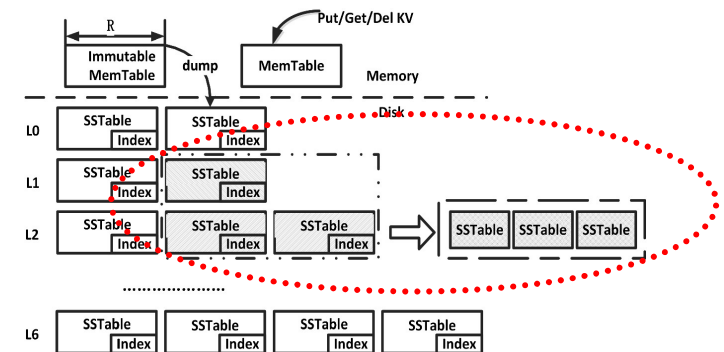
- ✓ SSTable (Sorted String Table)
- ✓ SSTable that exist in levels other than L0 have an **ordered state**
- ✓ Inside, there are data blocks, index blocks, bloom filter blocks, etc
- ✓ Index block
  - The data block containing the range containing the lookup key is used for lookup
  - Has a binary search data structure
- ✓ Bloom filter
  - Given an arbitrary key, this bit array may be used to determine if the key may exist or does not exist in the key set



# RocksDB Festival: RocksDB

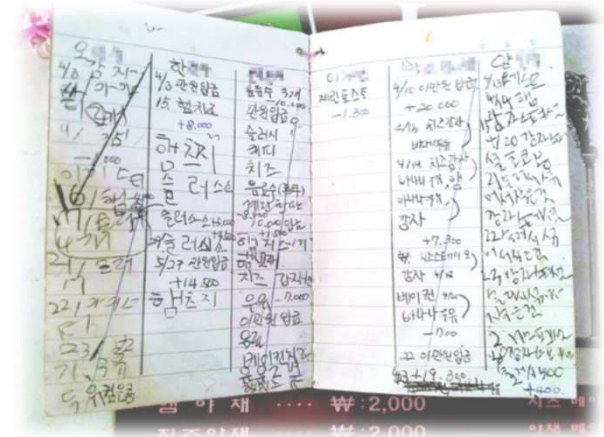
## ■ Main Terminology for RocksDB: Compaction

- ✓ Compaction is for update
- ✓ Prevents overlapping key-value pair from accumulating for existing key
- ✓ Data in L0 is the **hot** data, and the lower level is the **cold** data
- ✓ Compaction uses a **background thread**
- ✓ Various Type
  - Leveled compaction, Tiered(Universal) compaction, FIFO compaction



# RocksDB Festival: RocksDB

- Main Terminology for RocksDB: WAL(Write-Ahead-Log)
  - ✓ Record the WAL of memtable and disk for **every update**
  - ✓ Used for **recovery** in the event of an unexpected shutdown or error
  - ✓ Memtable securely as storage(disk) when flushed, the WAL is deleted.



# RocksDB Festival: RocksDB

## ■ RocksDB preview

```
root@linux-server-93:/home/choigunhee/hojin/RocksDB_Explorer# ls
appveyor.yml      db_stress_tool      issue_template.md   parsing_csv          third-party
AUTHORS           DEFAULT_OPTIONS_HISTORY.md  java                plugin              thirdparty.inc
buckifier         defs.bz1            LANGUAGE-BINDINGS.md PLUGINS.md          tools
build_tools      docs                librocksdb_debug.a  port               trace_replay
cache            DUMP_FORMAT.md     LICENSE.Apache      python_parser       USERS.md
cmake            env                LICENSE.leveldb     README.md          util
CMakeLists.txt   examples           logging             RocksDB_explorer_sh utilities
CODE_OF_CONDUCT.md file               make_config.mk      result_txt          Vagrantfile
CONTRIBUTING.md fuzz              Makefile            ROCKSDB_LITE.md    WINDOWS_PORT.md
COPYING          hdfs              memory              src.mk
coverage         HISTORY.md         memtable            table
db               include            monitoring           TARGETS
db_bench         INSTALL.md         options             test_util
```

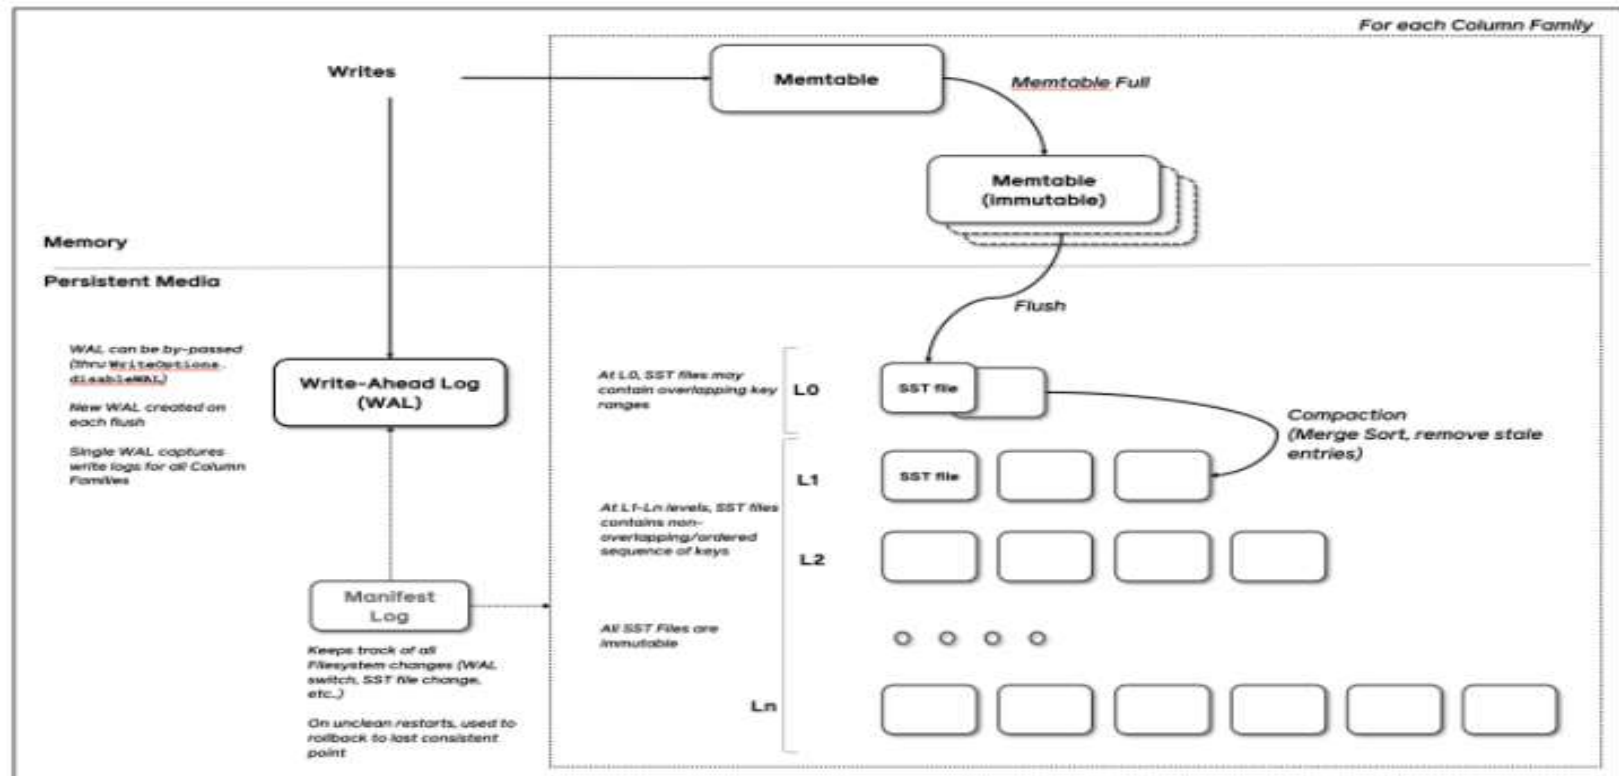
```
-rw-r--r-- 1 root root 37922501 7월 1 15:44 000214.sst
-rw-r--r-- 1 root root 37920200 7월 1 15:44 000216.sst
-rw-r--r-- 1 root root 37910828 7월 1 15:44 000219.sst
-rw-r--r-- 1 root root 37906740 7월 1 15:44 000221.sst
-rw-r--r-- 1 root root 37905482 7월 1 15:44 000224.sst
-rw-r--r-- 1 root root 37909294 7월 1 15:44 000227.sst
-rw-r--r-- 1 root root 66217640 7월 1 15:44 000228.log
-rw-r--r-- 1 root root 37892964 7월 1 15:44 000229.sst
-rw-r--r-- 1 root root 18621323 7월 1 15:44 000231.log
-rw-r--r-- 1 root root 16 7월 1 15:43 CURRENT
-rw-r--r-- 1 root root 37 7월 1 15:43 IDENTITY
-rw-r--r-- 1 root root 0 7월 1 15:43 LOCK
-rw-r--r-- 1 root root 572956 7월 1 15:44 LOG
-rw-r--r-- 1 root root 14592 7월 1 15:44 MANIFEST-000004
-rw-r--r-- 1 root root 6180 7월 1 15:43 OPTIONS-000007
```



# RocksDB Festival: RocksDB

## ■ Preview next week

- ✓ RocksDB Architecture
- ✓ RocksDB Operation (Compaction, WAL, Read, Write ...)
- ✓ RocksDB Benchmark



# Discussion

---

